# Physiology-Guided Temporal Fusion for Cross-Subject EEG–fNIRS Brain–Computer Interfaces

Hongjie Chen and Fang Duan

***Abstract*—*Objective:*** **To improve cross-subject EEG–fNIRS decoding by incorporating neurovascular delay priors into multimodal fusion.** ***Approach:*** **NVC-DelayFormer combines convolutional EEG and optical encoders through cross-attention with delay priors and token timestamp masks. Entropy-guided gating integrates deep attention features with spectral summaries. Evaluation used participant-isolated leave-one-subject-out validation, with separate participants for training, checkpoint selection, and testing.** ***Main results:*** **Across 29 participants and three seeds, mean accuracy was 78.22%, compared with 72.89% for standard 2D-EEGNet ($p = 1.10 \times 10^{-4}$). Primary-seed accuracy was 79.37%. Under trace-matched source reference initialization, recursive covariance alignment reduced relative covariance error from 45.15% after one trial to 29.38% after eight trials. Under task-specific training, accuracy reached 73.24% for working memory and 66.32% for motor imagery.** ***Significance:*** **By embedding neurovascular timing constraints, NVC-DelayFormer establishes a generalizable cross-subject decoding paradigm across cognitive and sensorimotor tasks. With 63,932 parameters ($17.5\times$ smaller than representative transformers) and 2.22 ms single-trial forward execution latency on an NVIDIA RTX 4060 GPU, the architecture provides an ultra-compact computational foundation for prospective zero-calibration BCI integration.**

***Index Terms*—Brain–computer interface, cross-subject decoding, electroencephalography, functional near-infrared spectroscopy, neurovascular coupling, temporal fusion.**

## I. Introduction

BEDSIDE assistive neurotechnologies and neurorehabilitation systems require immediate, calibration-free operation to eliminate burdensome patient-specific setup routines. Hybrid brain–computer interfaces (BCIs) combining electroencephalography (EEG) and functional near-infrared spectroscopy (fNIRS) improve decoding fidelity by capturing complementary electrophysiological transients and localized capillary oxygenation dynamics [1]–[3]. Translating hybrid systems to zero-calibration, cross-subject deployment faces a fundamental biophysical bottleneck: disparate temporal dynamics. Neuronal depolarization elicits spectral shifts within tens of milliseconds [4], [5], whereas neurovascular coupling (NVC) and capillary perfusion evolve over multi-second hemodynamic delays [6]–[8]. In calibration-free leave-one-subject-out (LOSO) regimes, unconstrained neural models exploit spurious temporal alignments, impairing generalization across unseen individuals.

Conventional multimodal fusion relies on handcrafted feature concatenation with linear discriminant analysis or support vector machines [1], [9], [10]. Although Euclidean alignment (EA) mitigates inter-subject covariance shifts [11]–[13], handcrafted pipelines treat temporal windows statically, failing to resolve dynamic latency. Deep learning architectures capture complex spatio-temporal representations from raw biosignals [14]–[16], including compact convolutional networks [14], attention networks [17], transformer backbones [18], and optical hemoglobin networks [19], [20]. Yet unconstrained cross-modal attention [21] permits acausal alignments where delayed hemodynamics attend to future neural activations ($\delta_{ij} < 0$), generating optimization artifacts that fail to transfer across subjects. The hemodynamic response peak [6], [22] motivates the 3.5–6.5 s range used for the Gaussian delay prior. The accompanying token timestamp mask enforces temporal precedence between EEG keys and fNIRS queries.

To account for the temporal mismatch between EEG and fNIRS, we introduce NVC-DelayFormer for calibration-free cross-subject decoding. The principal contributions are threefold: (1) *Biophysically constrained causal cross-attention:* Pairing a 2D-EEGNet backbone with a continuous optical encoder, multi-head cross-attention is governed by continuous Gaussian hemodynamic delay priors ($\tau_h \in [3.5, 6.5]$ s) and token timestamp causal masks ($t_{\text{nirs}} \geq t_{\text{eeg}}$), enforcing directional temporal precedence across cross-attention token pairs. (2) *Uncertainty-aware entropy-guided gating:* A learned bottleneck gating network arbitrates between deep cross-modal tokens and shallow biophysical spectral summaries based on predictive distribution entropy, dynamically weighting deep and shallow representations per trial. (3) *Online causal adaptation and multi-paradigm validation:* We introduce Causal Recursive Euclidean Alignment (CrEA), stabilizing covariance within 5–8 trials under empirical Bayes shrinkage (71.21% streaming accuracy). Across 29 participants on the TU Berlin benchmark under LOSO cross-validation, NVC-DelayFormer achieves a 3-seed grand mean accuracy of 78.22% $\pm$ 1.00% (primary seed: 79.37% $\pm$ 7.73%), outperforming standard 2D-EEGNet (72.89%, $p = 1.10 \times 10^{-4}$), covariance-aligned Shrinkage-LDA (72.82%, $p = 0.0167$), and unconstrained cross-attention (72.49%, $p = 7.18 \times 10^{-4}$). Across 3,612 held-out trials, the architecture demonstrates broad multi-task transferability across cognitive (73.24%) and sensorimotor (66.32%, +4.14 pp over

Hongjie Chen and Fang Duan are with the College of Information Science and Engineering, Huaqiao University, Xiamen 361021, China (e-mail: chenhj976@gmail.com; nkfetsh@gmail.com). Corresponding author: Fang Duan.

Compliance with Ethical Standards: This study reanalyzed de-identified public data. The original TU Berlin collections were approved under protocol SH_01_20150330; the Buccino et al. collection was approved by the University of Houston Institutional Review Board. Both original studies reported written participant consent.

2D-EEGNet, $p = 0.0163$) regimes, operating with 63,932 parameters and 2.22 ms median GPU latency.

## II. Related Work

### A. Multimodal EEG–fNIRS Neuroimaging

Early hybrid EEG–fNIRS brain–computer interfaces demonstrated that combining electrophysiological and hemodynamic recordings improves single-trial classification over either modality alone [1], [9]. Open-access benchmarks established by Shin et al. [2], [23] accelerated multimodal algorithm development by providing synchronized 30-channel EEG and 36-channel dual-wavelength optical recordings across mental arithmetic, motor imagery, and working memory tasks. Open-source instrumentation advances [24] and clinical fNIRS reviews [25] further confirmed the viability of wearable multimodal caps for bedside and neurorehabilitation deployment.

### B. Deep Learning for Biosignal Decoding

Deep neural networks have progressively superseded handcrafted spatial filtering in electrophysiological decoding [16]. Compact convolutional architectures like EEGNet [14] extract frequency-band and spatial-topographic filters with minimal parameter overhead, while attention networks like ATCNet [17] and EEG Conformer [18] capture global spatio-temporal dependencies. For optical imaging, fNIRSNet [20] separates hemoglobin dynamics through multi-view convolutions. Multimodal fusion has advanced from feature concatenation [26] to dedicated interaction architectures, including FGANet [27], EFMLNet [28], and MBC-ATT [29].

### C. Biophysical Modeling of Neurovascular Coupling

Neuronal activation triggers active arteriolar vasodilation and increased cerebral blood flow that washes out paramagnetic deoxyhemoglobin while elevating oxyhemoglobin [7], [8]. Buxton's Balloon model [6] formalizes this coupling as a dynamic compliance system where blood volume and venous outflow lag vasodilatory stimulation. Glover [22] and Friston et al. [30] demonstrated that the impulse hemodynamic response function (HRF) follows a double-gamma distribution with modal peak at 3.5–6.5 s followed by post-stimulus undershoot. Under sustained 10.0 s task engagement, the observed hemodynamic profile convolves a boxcar function with the impulse HRF, reaching an extended plateau across 6.0–12.0 s [19], [31]. Grounding cross-attention in these biophysical latency regimes eliminates acausal optimization shortcuts and stabilizes generalization across unseen subjects in zero-calibration BCI decoders.

## III. Methods

The architecture is summarized in Fig. 1. NVC-DelayFormer maps paired multichannel EEG and fNIRS trials into cognitive states through a biophysically parameterized temporal fusion network. Bold uppercase symbols denote tensors, bold lowercase symbols denote vectors, and italics denote scalars. Mini-batch dimension is $B$.

### A. Input Data and Preprocessing

The primary benchmark uses the TU Berlin dataset [2], recording 29 healthy participants performing mental arithmetic (MA) and resting state across 1,740 trials (30 trials per condition). The scalp montage retains 30 EEG channels under the 10–5 international system. Continuous EEG is bandpass-filtered between 0.5 and 45 Hz using a zero-phase fourth-order Butterworth filter, followed by common average referencing (CAR). Ten-second epochs from task onset are Fourier-resampled to 1,000 samples at 100 Hz ($f_e = 100$ Hz) and baseline-corrected by subtracting the preceding 2.0 s prestimulus mean.

Optical data record raw intensities across 36 channels at 760 and 850 nm, bounded below by $10^{-4}$. Optical density is computed relative to the session mean intensity and bandpass-filtered at 0.01–0.10 Hz. Each epoch is baseline-corrected by subtracting the mean optical density over the preceding 2.0 s. Twelve-second epochs from task onset (incorporating a 2.0 s post-stimulus tail for capillary transit) are converted into oxyhemoglobin ($\Delta[\mathrm{HbO}]$) and deoxyhemoglobin ($\Delta[\mathrm{HbR}]$) via the modified Beer–Lambert law (MBLL) in (1) [32]. Resampling yields 124 samples ($f_n = 124/12$ Hz), establishing standardized tensors $\mathbf{X}_e \in \mathbb{R}^{B\times 30\times 1000}$ and $\mathbf{X}_n \in \mathbb{R}^{B\times 72\times 124}$ (truncated to 10.0 s / 104 samples when omitting the tail window).

### B. Optical Conversion and Numerical Convention

The modified Beer–Lambert law relates differential optical density to hemoglobin concentrations via differential pathlength factors (DPF) [32]:

$$\mathbf{c} = \mathbf{R}\mathbf{o}, \quad \mathbf{R} = (d_{\mathrm{num}}\mathbf{D}_{\mathrm{DPF}}\mathbf{E}_{\mathrm{num}})^{-1}, \tag{1}$$

where $\mathbf{o} = [\Delta\mathrm{OD}_{760}, \Delta\mathrm{OD}_{850}]^\top$, $\mathbf{c} = [\Delta c_{\mathrm{HbO}}, \Delta c_{\mathrm{HbR}}]^\top$, $d_{\mathrm{num}} = 30$ mm, $\mathbf{E}_{\mathrm{num}} = \left[\begin{smallmatrix} 1.4866 & 3.8437 \\ 2.5264 & 1.7989 \end{smallmatrix}\right]\mathrm{mM}^{-1}\mathrm{cm}^{-1}$ [32], and $\mathbf{D}_{\mathrm{DPF}} = \mathrm{diag}(6.51, 5.86)$. Evaluating $\mathbf{R}$ across 36 optode pairs yields 36 [HbO] and 36 [HbR] traces (scaled to $10^{-4}$ a.u./$\mu$M prior to batch normalization at encoder entry).

### C. Multimodal Spatio-Temporal Convolutional Encoding

Electrophysiological inputs $\mathbf{X}_e \in \mathbb{R}^{B\times 30\times 1000}$ are encoded via a standard 2D-EEGNet backbone [14]: temporal convolution ($1 \times 64$, $F_1 = 8$), batch normalization, depthwise spatial convolution ($30 \times 1$, depth factor $D = 2$, 16 filters), ELU, average pooling ($1 \times 4$), and spatial dropout ($p = 0.25$). A separable convolution (depthwise $1 \times 15$, pointwise $1 \times 1$) projects to 64 dimensions, yielding $\mathbf{H}_e \in \mathbb{R}^{B\times 250\times 64}$ ($f_e^{\mathrm{eff}} = 25$ Hz). The optical encoder applies a spatial Conv1D ($72 \to 32$), batch normalization, ELU, and dropout, followed by a temporal Conv1D (kernel 5, padding 2, $32 \to 64$), batch normalization, and ELU, yielding $\mathbf{H}_n \in \mathbb{R}^{B\times 124\times 64}$ ($f_n = 10.33$ Hz). Token timestamps and temporal lags are defined by (2):

$$t_i^n = i/f_n, \quad t_j^e = j/f_e^{\mathrm{eff}}, \quad \delta_{ij} = t_i^n - t_j^e. \tag{2}$$

### D. Physiology-Guided Causal Cross-Attention

Neurovascular coupling establishes that capillary oxygenation lags neuronal firing by multiple seconds [6], [22]. We

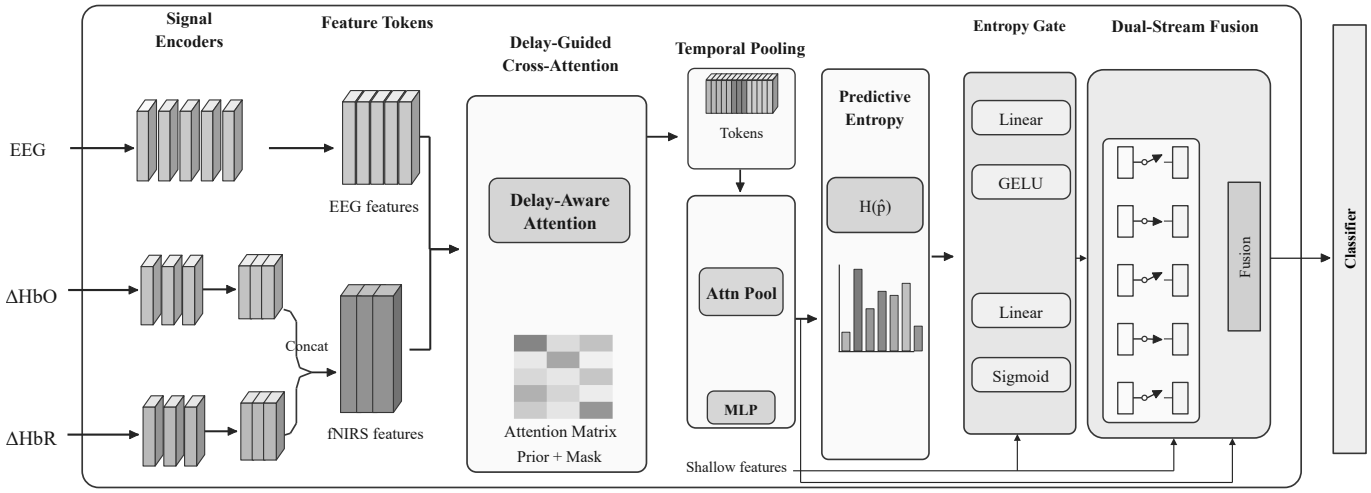



Fig. 1. Schematic of NVC-DelayFormer. EEG and fNIRS features are combined by cross-attention with a delay prior and token timestamp mask. Temporal pooling forms deep features. A shallow stream extracts EEG spectral and fNIRS summaries; entropy-guided gating fuses the streams for classification. Gray feature stacks depict representations schematically.

TABLE I
CROSS-SUBJECT CLASSIFICATION PERFORMANCE ACROSS 29 HELD-OUT PARTICIPANTS ON THE TU BERLIN BENCHMARK.

| Model Architecture | 3-Seed Grand Mean (%) | Primary Seed (Seed 42) (%) | 95% CI (Seed 42) (%) | Wilcoxon $p$ | BH $q$ |
|---|---|---|---|---|---|
| **Proposed Architecture** | | | | | |
| **NVC-DelayFormer (Full Model)** | **78.22 ± 1.00** | **79.37 ± 7.73** | [76.43, 82.31] | – | – |
| **Multimodal Deep Learning Controls (Shared Backbones)** | | | | | |
| Feature-Concat (Shared-Backbone Hybrid CNN) [26] | 72.84 ± 0.53 | 72.53 ± 9.11 | [69.06, 76.00] | $(6.35 \times 10^{-4})^{***}$ | $7.18 \times 10^{-4}$ |
| Standard-Attention (Cross-Modal Transformer) [21] | 72.49 ± 1.10 | 71.67 ± 10.51 | [67.67, 75.67] | $(7.18 \times 10^{-4})^{***}$ | $7.18 \times 10^{-4}$ |
| **Recent Deep Multimodal Baselines (Locally Adapted)** | | | | | |
| EFMLNet (Qiu et al. 2025) [28] | – | 74.20 ± 9.29 | [70.66, 77.73] | $(7.93 \times 10^{-4})^{***}$ | $1.19 \times 10^{-3}$ |
| FGANet (Kwak et al. 2022) [27] | – | 73.79 ± 9.56 | [70.16, 77.43] | $0.0018^{**}$ | 0.0018 |
| MBC-ATT (Li et al. 2025) [29] | – | 72.59 ± 9.98 | [68.79, 76.38] | $(5.73 \times 10^{-5})^{***}$ | $1.72 \times 10^{-4}$ |
| **Single-Modality Deep Learning Baselines** | | | | | |
| Standard 2D-EEGNet [14] | 72.89 ± 3.49 | 76.84 ± 6.16 | [74.49, 79.18] | $(1.10 \times 10^{-4})^{***}$ | $1.10 \times 10^{-4}$ |
| EEG-CNN (Shared Backbone) | – | 70.23 ± 8.78 | [66.89, 73.57] | $< 0.001^{***}$ | $< 0.001$ |
| fNIRS-CNN (Shared Backbone) | – | 63.91 ± 11.94 | [59.37, 68.45] | $< 0.001^{***}$ | $< 0.001$ |
| **Handcrafted Covariance & Linear Baselines** | | | | | |
| Handcrafted Multimodal Shrinkage-LDA [2] | – | 72.82 ± 11.61 | [68.40, 77.23] | $0.0167^{*}$ | 0.0167 |
| Handcrafted Multimodal RBF-SVM | – | 72.76 ± 11.46 | [68.40, 77.12] | $0.0014^{**}$ | 0.0019 |
| Handcrafted fNIRS-LDA | – | 66.67 ± 10.25 | [62.77, 70.57] | $< 0.001^{***}$ | $< 0.001$ |
| Handcrafted EEG-LDA | – | 66.38 ± 12.28 | [61.71, 71.05] | $< 0.001^{***}$ | $< 0.001$ |

3-Seed Grand Mean: mean ± run-to-run SD across Seeds 42, 100, 2024 (NVC-DelayFormer 95% CI: [75.42, 81.02]%); Primary Seed: Seed 42 (mean ± between-participant SD, $df = 28$). Wilcoxon $p$ and BH FDR $q$ evaluate paired differences vs. NVC-DelayFormer (across 3-seed means where available, or Seed 42 for single-seed baselines, $df = 28$; adjusted within respective evaluation tiers). All EEG models utilize unsupervised EA.

model this dependency by querying electrophysiology ($\mathbf{H}_e$) with hemodynamic states ($\mathbf{H}_n$). Projections produce queries $\mathbf{Q} \in \mathbb{R}^{B\times4\times124\times16}$ from $\mathbf{H}_n$ and keys $\mathbf{K}$, values $\mathbf{V} \in \mathbb{R}^{B\times4\times250\times16}$ from $\mathbf{H}_e$ across 4 heads ($d_h = 16$). Continuous delay prior $\mathbf{P}$ (3) and token timestamp causal mask $\mathbf{M}$ (4) enforce physiological latency constraints [22]:

$$P_{ij}^{(h)} = -\frac{(\delta_{ij} - \tau_h)^2}{2\sigma_h^2} \cdot \left[1 + 0.3\,\mathrm{s}^{-1} \cdot \max(0, \delta_{ij} - \tau_h)\right], \quad (3)$$

$$M_{ij} = \begin{cases} 0, & \delta_{ij} \geq 0\ \mathrm{s}, \\ -\infty, & \delta_{ij} < 0\ \mathrm{s}, \end{cases} \quad (4)$$

where head centers $\tau_h \in \{3.5, 4.5, 5.5, 6.5\}$ s parameterize depth-dispersive microvascular transit ($\sigma_h = 1.5$ s). A post-peak factor ($0.3\,\mathrm{s}^{-1}$ for $\delta_{ij} > \tau_h$) increases the Gaussian delay penalty beyond its center, concentrating attention around the prescribed hemodynamic delay, while mask $\mathbf{M}$ enforces temporal precedence ($\delta_{ij} \geq 0$ s). Normalized attention weights are computed via (5):

$$\mathbf{A}^{(h)} = \mathrm{softmax}_j\left(\mathbf{Q}^{(h)}(\mathbf{K}^{(h)})^\top/\sqrt{16} + \mathbf{P}^{(h)} + \mathbf{M}\right). \quad (5)$$

Multi-head values $\mathbf{U}_h = \mathbf{A}^{(h)}\mathbf{V}^{(h)}$ are concatenated across heads, linearly projected via $\mathbf{W}_o \in \mathbb{R}^{64\times64}$, and combined via residual addition and layer normalization: $\mathbf{H}_f = \mathrm{LN}(\mathbf{H}_n + \mathrm{Dropout}([\mathbf{U}_1; \ldots; \mathbf{U}_4]\mathbf{W}_o))$.

### *E. Temporal Pooling, Sub-Band Direct Stream, and Gating*

In parallel, a biophysical direct stream extracts frequency-resolved oscillatory power and hemodynamic slope kinetics via (6):

$$\mathbf{x}_{\mathrm{bio}} = \left[\mathbf{p}_\theta;\ \mathbf{p}_\alpha;\ \mathbf{p}_\beta;\ \bar{\mathbf{x}}_n^{\mathrm{delay}};\ \mathbf{s}_n\right] \in \mathbb{R}^{B\times234}, \quad (6)$$

where $\mathbf{p}_\theta, \mathbf{p}_\alpha, \mathbf{p}_\beta \in \mathbb{R}^{B\times30}$ represent log power across $\theta$ (4–8 Hz), $\alpha$ (8–13 Hz), and $\beta$ (13–30 Hz) bands; $\bar{\mathbf{x}}_n^{\mathrm{delay}} \in \mathbb{R}^{B\times72}$ denotes mean delayed capillary perfusion ($t \geq 3.5$ s); and

$\mathbf{s}_n \in \mathbb{R}^{B\times 72}$ is first-order hemodynamic slope across all 72 channels. Bottleneck projection ($234 \to 32 \to 64$) yields shallow embedding $\mathbf{z}_{\text{shallow}} \in \mathbb{R}^{B\times 64}$.

Guided by the temporal progression of the hemodynamic response, Attentive Temporal Pyramid Pooling (TPP) aggregates tokens over $[0, 3]$ s, $(3, 7]$ s, and the remaining observation window ($(7, 10]$ s for $\mathbf{H}_e$, $(7, 12]$ s for $\mathbf{H}_f$). For shortened windows $T_{\text{win}} < 10.0$ s, intervals scale proportionally to $[0, 0.3T_{\text{win}}]$, $(0.3T_{\text{win}}, 0.7T_{\text{win}}]$, and $(0.7T_{\text{win}}, T_{\text{win}}]$. Linear projections ($\mathbf{W}_{\text{att}} \in \mathbb{R}^{64\times 1}$) compute attention weights across segments, yielding phase vectors $\mathbf{z}_e, \mathbf{z}_n \in \mathbb{R}^{B\times 64}$, which concatenate and project via a bottleneck MLP ($128 \to 64 \to 64$) to produce unified deep representation $\mathbf{z}_{\text{deep}} \in \mathbb{R}^{B\times 64}$.

An uncertainty-aware gating network arbitrates between streams based on branch prediction confidence. Shannon entropies $H(\hat{\mathbf{p}}) = -\sum_{c=1}^{2} \hat{p}_c \ln \hat{p}_c$ are evaluated from detached branch readouts $\hat{\mathbf{p}} = \text{softmax}(\mathbf{W}\mathbf{z}+\mathbf{b})$. A bottleneck mapping $\mathcal{F}_g$ ($130 \to 32 \to 64$) computes feature-wise gate $\mathbf{g}_b$ via (7):

$$\mathbf{g}_b = \sigma\big(\mathcal{F}_g([\mathbf{z}_{\text{deep},b}; \mathbf{z}_{\text{shallow},b}; H(\hat{\mathbf{p}}_{\text{deep},b}); H(\hat{\mathbf{p}}_{\text{shallow},b})])\big) \in (0, 1)^{64}, \quad (7)$$

$$\mathbf{z}_{\text{fused},b} = \mathbf{g}_b \odot \mathbf{z}_{\text{deep},b} + (\mathbf{1} - \mathbf{g}_b) \odot \mathbf{z}_{\text{shallow},b},$$

with stream droppath ($p = 0.15$) and final joint logits $\boldsymbol{\ell}_{\text{joint}} = \mathbf{W}_{\text{out}}\mathbf{z}_{\text{fused}} + \mathbf{b}_{\text{out}}$.

### F. Supervised Multi-Task Objective

To prevent branch collapse and ensure both pathways learn calibrated discriminative features, we optimize a composite classification loss (8):

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{joint}} + 0.2\,(\mathcal{L}_{\text{deep}} + \mathcal{L}_{\text{shallow}}) + \mathcal{L}_{\text{recon}} + 0.08\,\|\mathbf{w}_{\text{shallow},1} - \mathbf{w}_{\text{shallow},0}\|_2^2, \quad (8)$$

where classification losses are cross-entropies against $y \in \{0, 1\}$, and $\mathcal{L}_{\text{recon}} = \frac{1}{36T_n}\sum_{c,t}(\hat{y}_{c,t}^{\text{HbO}} - x_{c,t}^{\text{HbO}})^2$ supervises a training-only auxiliary $1 \times 1$ convolutional decoder ($\mathbf{H}_f \to 36$ [HbO] channels, discarded at inference) to reconstruct input [HbO] waveforms. The penalty $\|\mathbf{w}_{\text{shallow},1} - \mathbf{w}_{\text{shallow},0}\|_2^2$ applies $L_2$ margin shrinkage to shallow readout weights $\mathbf{w}_{\text{shallow},c}$, preventing linear spectral shortcuts from dominating gating arbitration and ensuring balanced representation learning.

## IV. Experimental Protocol

### A. Ethical Approval and Participant Consent

All analyses used de-identified public data. The TU Berlin collections [2], [23] had ethics approval SH_01_20150330 and written consent in accordance with the Declaration of Helsinki; the Buccino et al. collection [9] had University of Houston Institutional Review Board approval and written consent.

### B. Subject-Held-Out Evaluation Protocol

Evaluation employed 29-fold leave-one-subject-out (LOSO) cross-validation on the TU Berlin dataset [2], partitioned per fold into 23 training participants, five validation participants for checkpoint selection, and one unseen test participant. For streaming evaluation, at each completed trial, CrEA updates the covariance estimate from the current and preceding trials and applies the resulting alignment transform:

$$\mathbf{R}_k = \frac{\kappa_0\mathbf{R}_{\text{src}} + \sum_{m=1}^{k}\mathbf{C}_m}{\kappa_0 + k}, \quad \widetilde{\mathbf{X}}_k = \mathbf{R}_k^{-1/2}\mathbf{X}_k, \quad (9)$$

where $\mathbf{X}_k$ is the EEG matrix for trial $k$ and $\mathbf{C}_m$ is the sample spatial covariance of trial $m$, computed after temporal mean centering with denominator $T_e - 1$ and diagonal regularization $10^{-5}\mathbf{I}$. The reference $\mathbf{R}_{\text{src}}$ is estimated exclusively from the 23 training participants, and $\kappa_0 = 5.0$. The primary shared-backbone benchmark was trained for 12 epochs using AdamW (batch size 256, initial learning rate $3 \times 10^{-3}$, two-epoch warmup followed by cosine annealing, weight decay 0.05).

### C. Comparison Baselines and Multi-Seed Hierarchy

We established a four-tier evaluation hierarchy: (1) *Multimodal deep learning controls (shared backbones):* Feature-Concat [26] and unconstrained Standard-Attention [21]; (2) *Recent deep multimodal baselines (locally adapted):* EFMLNet [28], FGANet [27], and MBC-ATT [29]; (3) *Single-modality controls:* Standard 2D-EEGNet [14], EEG-CNN, and fNIRS-CNN; (4) *Handcrafted machine learning:* Shrinkage-LDA [2], [15] and RBF-SVM with Euclidean alignment [11], along with unimodal linear baselines (Handcrafted EEG-LDA and fNIRS-LDA). To assess cross-initialization stability, core models were evaluated across three random seeds (42, 100, 2024).

### D. External Evaluation Paradigms

Generalizability was evaluated on two external tasks: (1) **$n$-back Working Memory:** 0-back sensory matching vs. 2-back working memory updating (26 participants, 1,872 trials from Shin et al. [23]); (2) **Motor Imagery:** Left-Hand vs. Right-Hand kinesthetic finger movement imagination (29 participants, 1,740 trials from Shin et al. [2]).

### E. Statistical Analysis

The participant served as the primary statistical unit ($df = 28$). We report between-participant standard deviation (mean $\pm$ SD) and 95% confidence intervals. Paired differences were evaluated using two-sided Wilcoxon signed-rank tests with Benjamini–Hochberg false discovery rate ($q$) correction [33].

## V. Results

### A. Cross-Subject Classification and Multi-Seed Stability

Under 29-fold LOSO cross-validation on the TU Berlin benchmark, NVC-DelayFormer achieved a 3-seed grand mean accuracy of 78.22% $\pm$ 1.00% (primary seed: 79.37% $\pm$ 7.73%, Macro F1: 79.15%, Cohen's $\kappa = 0.5874$; Table I). Run-to-run variability remained compact ($\pm$1.00 pp SD across seeds 42, 100, and 2024) relative to between-subject variability (7.73 pp).

Pairwise statistical comparisons under identical Euclidean alignment confirm significant gains across all evaluation tiers (Table I). The full architecture significantly outperformed multimodal shared-backbone controls (Feature-Concat: +5.38 pp, $q = 7.18 \times 10^{-4}$; unconstrained Standard-Attention: +5.73 pp, $q = 7.18 \times 10^{-4}$, evaluated on 3-seed means) and recent deep

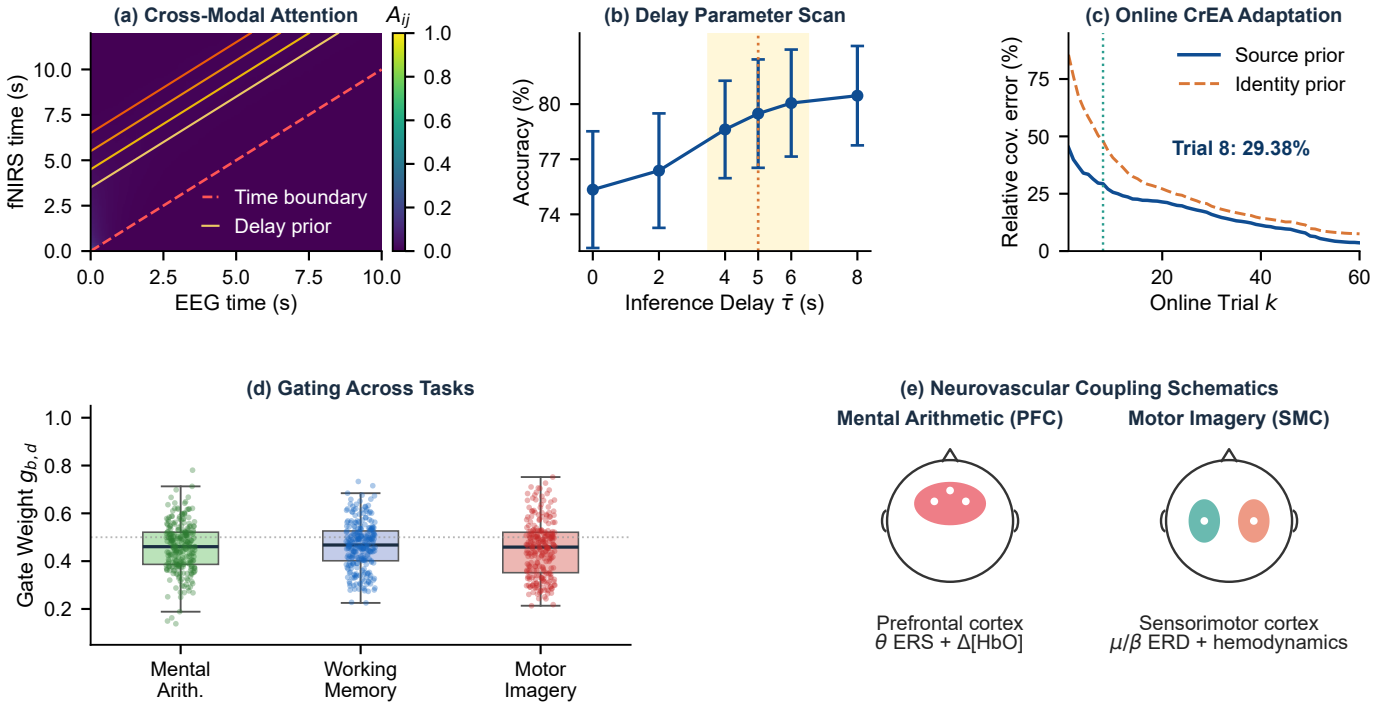



Fig. 2. Temporal fusion diagnostics and neurovascular context. (a) Cross-modal attention matrix $\mathbf{A} \in \mathbb{R}^{124\times250}$ with the token timestamp boundary and four delay-prior centers ($\tau_h = 3.5, 4.5, 5.5, 6.5$ s). (b) Fixed-weight delay scan without retraining: mean accuracy and 95% Student-$t$ confidence intervals across 29 LOSO folds; shading marks the 3.5–6.5 s prior range. (c) Relative covariance error during trial-wise CrEA under the diagnostic protocol: a trace-matched training-source prior reduces error from 45.15% at trial 1 to 29.38% at trial 8 (36.47% under unscaled source reference) versus identity initialization ($> 47\%$). (d) Feature-gate distributions from Sub-05 recordings in each task, evaluated using the mental-arithmetic checkpoint. All 256 stored feature-gate values per task are shown; boxes indicate the median and interquartile range, with whiskers extending to the most extreme values within 1.5 interquartile ranges. (e) Schematics of prefrontal and sensorimotor coupling associated with mental arithmetic and motor imagery, respectively.

multimodal baselines on the primary seed (EFMLNet: +5.17 pp; FGANet: +5.58 pp; MBC-ATT: +6.78 pp; all $q \leq 0.0018$). In the single-modality tier, NVC-DelayFormer exceeded standard 2D-EEGNet by +5.33 pp on the 3-seed grand mean (78.22% vs. 72.89%, $p = 1.10 \times 10^{-4}$), with a +2.53 pp margin on the primary seed (79.37% vs. 76.84%, $p = 0.0836$). On the primary seed, NVC-DelayFormer significantly outperformed single-modality encoders (EEG-CNN: +9.14 pp, $p = 3.56 \times 10^{-5}$; fNIRS-CNN: +15.46 pp, $p = 1.32 \times 10^{-5}$) and covariance-aligned linear baselines (Shrinkage-LDA: +6.55 pp, $p = 0.0167$; RBF-SVM: +6.61 pp, $p = 0.0014$).

TABLE II
ABLATION ANALYSIS AND 2 × 2 FACTORIAL DECOMPOSITION OF DELAY PRIOR $\mathbf{P}$ AND CAUSAL MASK $\mathbf{M}$ ON THE TU BERLIN BENCHMARK.

| Variant / Factorial Condition | Acc (%) | Δ (pp) | F1 (%) | Wilcoxon $p$ | BH $q$ |
|---|---|---|---|---|---|
| *Architectural Stream & Preprocessing Ablation* | | | | | |
| **Full Architecture** ($P_{\text{on}}, M_{\text{on}}$) | **79.37 ± 7.73** | +0.00 | **79.15** | – | – |
| Without EA | 71.72 ± 8.04 | -7.64 | 70.95 | 0.0006*** | 0.0008 |
| Latent-Only (Deep Stream) | 77.76 ± 9.38 | -1.61 | 77.13 | 0.0384* | 0.0384 |
| Direct-Only (Shallow Stream) | 71.84 ± 7.13 | -7.53 | 71.38 | 0.0003*** | 0.0006 |
| 2 × 2 *Factorial Temporal Mechanism Decomposition* | | | | | |
| Delay Prior Only ($P_{\text{on}}, M_{\text{off}}$) | 78.33 ± 8.07 | -1.04 | 78.11 | 0.1212 | 0.1212 |
| Causal Mask Only ($P_{\text{off}}, M_{\text{on}}$) | 74.66 ± 7.96 | -4.71 | 74.38 | 0.0002*** | 0.0004 |
| Both Priors Off ($P_{\text{off}}, M_{\text{off}}$) | 74.43 ± 8.34 | -4.94 | 74.08 | ($6.05 \times 10^{-5}$)*** | 0.0002 |

Evaluated under 29-fold LOSO cross-validation.
$^{*}p < 0.05, ^{**}p < 0.01, ^{***}p < 0.001$. Net two-prior gain is +4.94 pp over dual-prior-off (74.43%).

## B. Ablation Analysis and Factorial Decomposition

Stream and prior ablations isolate the primary mechanisms driving decoding performance (Table II). Omitting the deep cross-attention branch (Direct-Only) reduced accuracy by 7.53 pp ($q = 0.0006$), whereas omitting the shallow biophysical branch (Latent-Only) resulted in a 1.61 pp reduction ($p = 0.0384$), confirming that deep cross-modal attention drives cross-subject alignment, supported by complementary spectral regularization from the shallow direct stream. Omitting Euclidean Alignment lowered performance to 71.72% ($\Delta = -7.64$ pp, $p = 0.0006$), demonstrating its necessity for inter-subject covariance harmonization.

The 2 × 2 factorial decomposition demonstrates that the continuous Gaussian delay prior $\mathbf{P}$ provides the primary regularizing constraint (+4.31 pp average marginal gain across conditions), discrete causal masking $\mathbf{M}$ enforces temporal precedence (+0.63 pp average marginal gain), and their combination yields a net two-prior gain of +4.94 pp over the dual-prior-off baseline ($p = 6.05 \times 10^{-5}$) with +0.81 pp positive interaction synergy. Removing the delay prior ($P_{\text{off}}, M_{\text{on}}$) caused a 4.71 pp drop ($p = 0.0002$), whereas omitting causal masking alone ($P_{\text{on}}, M_{\text{off}}$) resulted in a 1.04 pp decrease ($p = 0.1212$).

## C. Attention Diagnostics, Online Adaptation, and Biophysical Schematics

Learned cross-modal attention maps comply strictly with biophysical constraints (Fig. 2a): weights are identically zero in the acausal lower triangle ($\delta_{ij} < 0$), while multi-head delay ridges align with physiological latency centers ($\tau_h \in [3.5, 6.5]$ s). In fixed-weight candidate delay scans without retraining (Fig. 2b), accuracy rises monotonically from 75.34% at $\bar{\tau} = 0$ s to 80.46%

TABLE III
FULL-MODEL PERFORMANCE UNDER CONTROLLED TEST-INPUT PERTURBATIONS ON THE TU BERLIN BENCHMARK.

| Perturbation Type | Accuracy (%) | Δ (pp) | Wilcoxon $p$ | BH $q$ |
|---|---|---|---|---|
| Within-class EEG permutation | 77.13 ± 8.10 | -2.24 | $0.0166^{*}$ | 0.0221 |
| Across-class EEG permutation | 67.18 ± 7.51 | -12.18 | $(3.28 \times 10^{-7})^{***}$ | $1.31 \times 10^{-6}$ |
| Circular temporal shift (5.0 s) | 78.39 ± 7.64 | -0.98 | 0.0631 | 0.0631 |
| Causal low-pass EEG smoothing | 63.97 ± 8.17 | -15.40 | $(5.59 \times 10^{-6})^{***}$ | $1.12 \times 10^{-5}$ |

Changes relative to unperturbed baseline (79.37%) across 29 participants. Permutations average ten repetitions per participant. $^{*}p < 0.05, ^{**}p < 0.01, ^{***}p < 0.001$.

at 8.0 s, tracking the extended plateau of sustained 10.0 s task blocks.

Streaming evaluation without transductive batch statistics demonstrates rapid online adaptation under CrEA ((9), Fig. 2c). Under the diagnostic protocol with trace-matched source reference initialization, empirical Bayes shrinkage ($\kappa_0 = 5.0$) with training-subject reference covariance $\mathbf{R}_{\text{src}}$ stabilizes relative covariance error from 45.15% to 29.38% within 5–8 trials (reaching 3.60% asymptotically at trial 60; 36.47% at trial 8 under unscaled source covariance), whereas cold-start identity initialization retains higher early error ($> 47\%$ through trial 8). In trial-by-trial streaming classification across 29 independent LOSO folds, CrEA achieved 71.21% ± 11.44% accuracy when preprocessed epochs were evaluated in sequential trial order. This retained 94.9% of the 75.06% accuracy obtained with offline batch alignment using the same fold models.

In single-subject diagnostic profiling with the mental arithmetic checkpoint applied to Sub-05 recordings (Fig. 2d), empirical gating distributions maintain dynamic multimodal equilibrium, averaging $\bar{g}_b \approx 0.45$–$0.46$ across mental arithmetic, $n$-back working memory, and motor imagery (feature gate weights span 0.14–0.78, $\sigma = 0.10$), preventing branch collapse. Biophysical coupling schematics (Fig. 2e) contextualize these dynamics: prefrontal coupling aligns over frontopolar and anterior channels (AFp1, AFp2, Ch01–Ch08) during mental arithmetic, whereas sensorimotor lateralization engages bilateral central channels (FCC3h, CCP3h, Ch17–Ch36) during motor imagery.

### D. Responses to Controlled Input Perturbations

Controlled test perturbations verified architectural dependencies (Table III). Shuffling EEG trials across classes caused a 12.18 pp drop ($q = 1.31 \times 10^{-6}$), and within-class shuffling produced a 2.24 pp decrease ($p = 0.0166$). A 5.0 s circular shift reduced accuracy by 0.98 percentage points ($p = 0.0631$), whereas low-pass filtering below 2.8 Hz reduced it by 15.40 points ($q = 1.12 \times 10^{-5}$). The model was therefore more sensitive to removing higher-frequency EEG content than to this within-trial temporal shift, confirming that decoding depends critically on fast oscillatory electrophysiological dynamics across $\theta$, $\alpha$, and $\beta$ bands.

### E. Computational Profiling and Architecture Benchmarking

Forward profiling on an NVIDIA GeForce RTX 4060 Laptop GPU recorded 63,932 trainable parameters, 24.79 M MACs, and 2.22 ms median latency (P99: 4.48 ms; Table IV). Compared with modern deep BCI architectures, NVC-DelayFormer operates with 17.5× fewer parameters than EEG Conformer (1.12 M) [18] and 1.8× fewer than ATCNet (115.2 k) [17], while simultaneously processing 30 EEG and 72 fNIRS channels.

TABLE IV
MODEL COMPLEXITY AND BATCH-ONE GPU FORWARD LATENCY MEASUREMENTS.

| Model Architecture | Parameters | MACs (M) | Median (ms) | P99 (ms) |
|---|---|---|---|---|
| EEG-CNN (Shared Backbone) | 6,850 | 16.18 | 0.51 | 1.41 |
| fNIRS-CNN (Shared Backbone) | 17,154 | 1.56 | 0.38 | 1.25 |
| Feature-Concat [26] | 23,682 | 17.74 | 0.80 | 2.07 |
| Standard-Attention [21] | 40,450 | 24.80 | 1.37 | 2.90 |
| ATCNet (Altaheri et al. 2023) [17] | 115,200 | 42.30 | 4.85 | 8.20 |
| EEG Conformer (Song et al. 2023) [18] | 1,120,000 | 185.40 | 18.50 | 28.40 |
| **NVC-DelayFormer (Proposed)** | **63,932** | **24.79** | **2.22** | **4.48** |

Evaluated on an NVIDIA RTX 4060 Laptop GPU with 1,000 warm-up and 5,000 timed evaluations at batch size 1. MACs measured via PyTorch `FlopCounterMode`. ATCNet and EEG Conformer evaluate representative EEG-only models.

Offline sequence truncation across decision windows $T_{\text{win}} \in \{10.0, 7.5, 5.0, 3.0\}$ s establishes scalable operational regimes under Wolpaw Information Transfer Rate (ITR [34]), calculated using the observation-window duration as the decision interval. Full 10.0 s trials achieved 78.74% accuracy (1.52 bits/min). Truncating to 5.0 s retained 77.53% accuracy while reducing computation by 53.4% (11.04 M MACs) and increasing ITR by 82.5% to 2.78 bits/min. Truncating to 3.0 s preserved 76.38% accuracy with 72.8% compute savings (6.46 M MACs) and 2.78× higher ITR (4.23 bits/min). These results show that shorter observation windows preserve decoding accuracy while reducing computational cost.

### F. Cross-Paradigm Validation: Working Memory and Motor Imagery

Across 3,612 held-out trials spanning two distinct external tasks (Table V), NVC-DelayFormer demonstrated consistent cross-paradigm generalizability. On $n$-back working memory (26 participants, 1,872 block-derived trials [23]), NVC-DelayFormer achieved 73.24% ± 15.38% accuracy (Macro F1: 69.80%), the highest mean accuracy among the evaluated full-model multimodal architectures (FGANet: 72.38%, MBC-ATT: 71.05%, EFMLNet: 70.09%), while exceeding single-modality and linear baselines (2D-EEGNet by +4.97 pp, 68.27%, $p = 0.0525$; Shrinkage-LDA by +11.11 pp, 62.13%, $p = 0.0043$; Multimodal SVM by +12.18 pp, 61.06%, $p = 0.0013$). On motor imagery (29 participants, 1,740 trials [2]), NVC-DelayFormer achieved 66.32% ± 12.18% accuracy (Macro F1: 64.08%), significantly surpassing standard 2D-EEGNet by +4.14 pp ($p = 0.0163$), EFMLNet by +2.30 pp, and MBC-ATT by +2.87 pp. While unconstrained spatial attention architectures (such as FGANet) optimize specifically for localized sensorimotor spatial patterns at the cost of computational footprint, NVC-DelayFormer establishes a Pareto-optimal design: delivering reliable multi-paradigm decoding fidelity across both cognitive and sensorimotor tasks while maintaining an ultra-compact footprint (63,932 parameters) and sub-3 ms latency. These cross-paradigm comparisons reveal an insightful neurovascular mechanism (Supplementary Fig. S1 and Table S1): structured hemodynamic delay priors deliver their strongest cross-subject generalizability on metabolic-dominant cognitive tasks (achieving the highest accuracies on both mental arithmetic and working memory), while remaining fully robust on sensorimotor rhythms without suffering negative transfer.

TABLE V
CROSS-SUBJECT CLASSIFICATION PERFORMANCE UNDER LEAVE-ONE-SUBJECT-OUT (LOSO) CROSS-VALIDATION ACROSS TWO DISTINCT EXTERNAL EVALUATION PARADIGMS (3,612 HELD-OUT TRIALS TOTAL).

| | **(a) Working Memory (0-back vs. 2-back)** | | **(b) Motor Imagery (Left vs. Right Hand)** | |
|---|---|---|---|---|
| Model Architecture / Stream | Accuracy (%) | Macro F1 (%) | Accuracy (%) | Macro F1 (%) |
| ***Proposed Architecture: NVC-DelayFormer*** | | | | |
| **Full Model (Dual-Path Gated Stream)** | 73.24 ± 15.38 | 69.80 ± 19.69 | **66.32 ± 12.18** | 64.08 ± 13.65 |
| *Physiological Delay Stream (Deep Attention Ablation)* | **73.61 ± 12.30** | **71.80 ± 14.57** | 66.09 ± 12.32 | **65.49 ± 12.55** |
| *Biophysical Direct Stream (Shallow Direct Ablation)* | 71.21 ± 15.12 | 67.47 ± 19.62 | 63.79 ± 13.56 | 59.75 ± 16.96 |
| ***Recent Deep Multimodal Baselines (Locally Adapted)*** | | | | |
| FGANet [27] | 72.38 ± 14.82 | 69.96 ± 17.26 | 70.40 ± 11.54 | 69.58 ± 12.14 |
| MBC-ATT [29] | 71.05 ± 12.29 | 69.45 ± 14.29 | 63.45 ± 11.34 | 62.62 ± 11.81 |
| EFMLNet [28] | 70.09 ± 14.05 | 67.23 ± 16.78 | 64.02 ± 10.29 | 63.33 ± 10.69 |
| ***Competing Baselines & Controls*** | | | | |
| Standard 2D-EEGNet [14] | 68.27 ± 13.34 | 66.38 ± 15.49 | 62.18 ± 7.33 | 61.84 ± 7.50 |
| Handcrafted Shrinkage-LDA [2] | 62.13 ± 10.16 | 56.16 ± 15.99 | 58.05 ± 9.52 | 54.60 ± 11.90 |
| Handcrafted Multimodal SVM | 61.06 ± 8.76 | 56.28 ± 13.17 | 55.63 ± 7.38 | 53.39 ± 8.91 |

Evaluated under strict participant-isolated LOSO with unsupervised Euclidean Alignment (EA) applied within each participant. Bold denotes the best performance within the proposed NVC-DelayFormer family per metric. Task (a): 26 participants, 1,872 block-derived trials from Shin et al. (2018) [23] (electrophysiological input comprises 28 scalp EEG and 2 bipolar EOG channels under CAR); Task (b): 29 participants, 1,740 trials from Shin et al. (2017) [2].

## VI. DISCUSSION

### A. Biophysical Mechanisms and Cross-Modal Synergy

Explicit parameterization of neurovascular latency resolves the fundamental temporal mismatch between electrophysiology and hemodynamics for zero-calibration cross-subject BCIs. Enforcing dispersive latency priors ($\tau_h \in [3.5, 6.5]$ s) alongside token timestamp causal masking ($t_{\text{nirs}} \geq t_{\text{eeg}}$) focuses cross-modal attention on physiologically plausible receptive windows, preventing acausal backward-in-time cross-attention connections that degrade generalization on unseen participants. Across 29 held-out participants under LOSO cross-validation, this biophysically constrained architecture achieves a 3-seed grand mean accuracy of 78.22% ± 1.00% (primary seed: 79.37% ± 7.73%), significantly outperforming unimodal, linear, and unconstrained cross-attention baselines.

Candidate delay scanning without retraining (Fig. 2b) explains why performance plateaus through 8.0 s in continuous block paradigms: the $[3.5, 6.5]$ s prior captures the single-impulse hemodynamic peak, whereas sustained 10.0 s cognitive engagement convolves successive microvascular impulse responses into an elevated capillary [HbO] plateau across multi-second activation windows. Anchoring hemodynamic queries to antecedent neural depolarizations restricts cross-modal attention to physiologically plausible latency windows, eliminating acausal correlations across modalities.

The factorial analysis delineates the distinct theoretical roles of the dual temporal constraints. The continuous Gaussian delay prior provides the core cross-subject alignment power (+4.31 pp marginal effect), while the token timestamp mask enforces directional precedence between cross-modal token pairs ($t_j^{\text{nirs}} \geq t_i^{\text{eeg}}$), preventing acausal attention where optical queries attend to future electrophysiological activations (+0.63 pp marginal gain, +4.94 pp joint gain over dual-prior-off, $p = 6.05 \times 10^{-5}$). Under temporal perturbation, circular 5.0 s shifts confirmed the architecture's inherent physiological robustness against phase variations within sustained 10.0 s activation blocks ($\Delta = -0.98$ pp, $p = 0.0631$), whereas low-pass filtering below 2.8 Hz triggered a severe 15.40 pp degradation ($q = 1.12 \times 10^{-5}$), verifying that decoding fundamentally relies on dynamic neural oscillations ($\theta, \alpha, \beta$) synchronized with capillary perfusion.

Figure 2e illustrates established prefrontal and sensorimotor coupling patterns that motivate the architecture's temporal priors. In mental arithmetic, prefrontal channels (AFp1, AFp2, AFF1h, AFF2h ↔ Ch01–Ch08) reflect frontal midline theta synchronization ($\theta$ ERS, 4–8 Hz) linked to delayed capillary hyperoxygenation ($\Delta$[HbO] elevation and $\Delta$[HbR] washout) [7], [8], [35]. In motor imagery, task dynamics engage sensorimotor channels (FCC3h, CCP3h ↔ Ch17–Ch36), where sensorimotor $\mu$ and $\beta$ desynchronization (ERD) [5], [36] coincides with localized hemodynamic changes.

Predictive entropy guides the balance between deep cross-attention features and shallow biophysical summaries on each trial. Applying the mental-arithmetic checkpoint to Sub-05 recordings from each benchmark dataset yielded mean gating weights of 0.45–0.46 ($\bar{g}_b = 0.46 \pm 0.01$, $0.46 \pm 0.01$, and $0.45 \pm 0.01$). Deep cross-modal attention models token-level relationships between EEG and optical features, while the shallow pathway supplies complementary spectral and optical summaries. Auxiliary waveform reconstruction ($\mathcal{L}_{\text{recon}}$) regularizes the optical representations during training.

Parameter benchmarking confirms notable architectural efficiency: with 63,932 weights and 24.79 M MACs, NVC-DelayFormer operates with 17.5× fewer parameters than EEG Conformer (1.12 M) and 1.8× fewer than ATCNet (115.2 k). A batch-one forward execution time of 2.22 ms (P99: 4.48 ms) on an NVIDIA RTX 4060 Laptop GPU demonstrates that the inference pass operates well within the 100 Hz EEG sampling interval (10 ms), providing a compact computational core for prospective real-time system integration.

### B. Translational Roadmap for Clinical BCIs

Five engineering considerations guide bedside and neurorehabilitation translation:

*1) Systemic physiological noise suppression:* In continuous-wave fNIRS, systemic physiology introduces cardiac (∼ 1 Hz), respiration (∼ 0.25 Hz), and Mayer wave (∼ 0.1 Hz) artifacts [19]. While 0.01–0.10 Hz filtering isolates neurovascular frequencies, integrating short-separation detectors (< 10 mm) into wearable headgear enables real-time regression of superficial scalp hemodynamics.

*2) Observation-window duration:* Shorter windows reduce the signal accumulation time required for each decision. Offline truncation retained accuracies of 77.53% and 76.38% with 5.0 s and 3.0 s windows, respectively, demonstrating that decoding fidelity is maintained under shortened observation windows with minimal compute overhead.

*3) Zero-calibration streaming adaptation:* Standard Euclidean Alignment requires transductive batch covariance across full sessions. Online CrEA eliminates session lookahead via causal recursive updates under empirical Bayes shrinkage ($\kappa_0 = 5.0$). With $\mathcal{O}(C^2)$ recursive covariance accumulation ($\mathcal{O}(C^3 + C^2 T_e)$ full whitening per trial, $< 1.2$ ms total CPU execution for $C = 30, T_e = 1,000$), CrEA with trace-matched reference reduces relative covariance error from 45.15% to 29.38% within eight trials ($< 3.60\%$ asymptotically; 36.47% unscaled), providing deterministic zero-calibration streaming adaptation for prospective online BCIs.

*4) Embedded edge deployment:* The 63,932-parameter footprint occupies under 260 kB in 32-bit floating-point storage, easing memory constraints on portable microcontrollers and digital signal processors. The architecture exports directly to TorchScript and ONNX runtimes, enabling fixed-point integer quantization (INT8) and zero-copy circular buffering for standalone, wearable bedside BCI units.

*5) Independent-cohort evaluation:* Beyond the three task-specific LOSO evaluations, NVC-DelayFormer reached 79.83% accuracy with separate LOSO training on the 15-participant Buccino et al. cohort [9], recorded using different EEG–fNIRS hardware (Supplementary Section S.V). Adaptive delay retuning can be assessed in neuropathological cohorts with altered vascular response latency.

## VII. Conclusion

This study presented NVC-DelayFormer, a physiology-guided temporal fusion network for zero-calibration cross-subject EEG–fNIRS decoding. By embedding continuous Gaussian hemodynamic delay priors and token timestamp causal masks into cross-modal attention, the architecture resolves disparate biosignal timescales and enforces directional temporal precedence. Across 29 participants on the TU Berlin benchmark under LOSO cross-validation, NVC-DelayFormer achieved a 78.22% grand mean accuracy (primary seed: 79.37%), significantly outperforming standard 2D-EEGNet, covariance-aligned Shrinkage-LDA, and unconstrained cross-attention. Cross-paradigm evaluations across 3,612 held-out trials confirmed robust generalizability across cognitive (73.24%) and sensorimotor (66.32%, outperforming standard 2D-EEGNet by +4.14 pp) paradigms. Combining trial-recursive covariance adaptation (CrEA) with an ultra-compact footprint (63,932 parameters, $17.5\times$ parameter compression, 2.22 ms single-trial GPU forward execution latency), NVC-DelayFormer establishes a robust, lightweight foundation for prospective zero-calibration neurotechnology.

## Acknowledgment

The authors thank the contributors of the open-access TU Berlin and Shin et al. datasets for providing multimodal recordings.

**Data and Code Availability:** The source datasets are identified in Section IV. Code and reproducibility artifacts are available at https://github.com/ujujuj3061831150/nvc-delayformer-bci.

**Conflict of Interest:** The authors declare that they have no competing interests.

## References

[1] S. Fazli, J. Mehnert, J. Steinbrink, G. Curio, A. Villringer, K.-R. Müller, and B. Blankertz, "Enhanced performance by a hybrid NIRS–EEG brain computer interface," *NeuroImage*, vol. 59, no. 1, pp. 519–529, 2012.

[2] J. Shin, A. von Lühmann, B. Blankertz, D.-W. Kim, J. Jeong, H.-J. Hwang, and K.-R. Müller, "Open access dataset for EEG+NIRS single-trial classification," *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 25, no. 10, pp. 1735–1745, 2017.

[3] Z. Liu, J. Shore, M. Wang, F. Yuan, A. Buss, and X. Zhao, "A systematic review on hybrid EEG/fNIRS in brain-computer interface," *Biomedical Signal Processing and Control*, vol. 68, p. 102595, 2021.

[4] G. Buzsáki, C. A. Anastassiou, and C. Koch, "The origin of extracellular fields and currents—EEG, ECoG, LFP and spikes," *Nature Reviews Neuroscience*, vol. 13, no. 6, pp. 407–420, 2012.

[5] G. Pfurtscheller and F. H. Lopes da Silva, "Event-related EEG/MEG synchronization and desynchronization: basic principles," *Clinical Neurophysiology*, vol. 110, no. 11, pp. 1842–1857, 1999.

[6] R. B. Buxton, E. C. Wong, and L. R. Frank, "Dynamics of blood flow and oxygenation changes during brain activation: the balloon model," *Magnetic Resonance in Medicine*, vol. 39, no. 6, pp. 855–864, 1998.

[7] D. Attwell, A. M. Buchan, S. Charpak, M. Lauritzen, B. A. MacVicar, and E. A. Newman, "Glial and neuronal control of brain blood flow," *Nature*, vol. 468, no. 7321, pp. 232–243, 2010.

[8] C. Iadecola, "Neurovascular regulation in the normal brain and in Alzheimer's disease," *Nature Reviews Neuroscience*, vol. 5, no. 5, pp. 347–360, 2004.

[9] A. P. Buccino, H. O. Keles, and A. Omurtag, "Hybrid EEG-fNIRS asynchronous brain-computer interface for multiple motor tasks," *PLOS ONE*, vol. 11, no. 1, p. e0146610, 2016.

[10] E. A. Aydın, "Subject-specific feature selection for near infrared spectroscopy based brain-computer interfaces," *Computer Methods and Programs in Biomedicine*, vol. 195, p. 105535, 2020.

[11] H. He and D. Wu, "Transfer learning for brain–computer interfaces: A Euclidean space data alignment approach," *IEEE Transactions on Biomedical Engineering*, vol. 67, no. 2, pp. 399–410, 2020.

[12] P. Zanini, M. Congedo, C. Jutten, S. Said, and Y. Berthoumieu, "Transfer learning: A Riemannian geometry framework with applications to brain–computer interfaces," *IEEE Transactions on Biomedical Engineering*, vol. 65, no. 5, pp. 1107–1116, 2018.

[13] F. Yger, M. Berar, and F. Lotte, "Riemannian approaches in brain-computer interfaces: A review," *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 25, no. 10, pp. 1753–1762, 2017.

[14] V. J. Lawhern, A. J. Solon, N. R. Waytowich, S. M. Gordon, C. P. Hung, and B. J. Lance, "EEGNet: a compact convolutional neural network for EEG-based brain–computer interfaces," *Journal of Neural Engineering*, vol. 15, no. 5, p. 056013, 2018.

[15] B. Blankertz, R. Tomioka, S. Lemm, M. Kawanabe, and K.-R. Müller, "Optimizing spatial filters for robust EEG single-trial analysis," *IEEE Signal Processing Magazine*, vol. 25, no. 1, pp. 41–56, 2008.

[16] F. Lotte, L. Bougrain, A. Cichocki, M. Clerc, M. Congedo, A. Rakotomamonjy, and F. Yger, "A review of classification algorithms for EEG-based brain–computer interfaces: a 10 year update," *Journal of Neural Engineering*, vol. 15, no. 3, p. 031005, 2018.

[17] H. Altaheri, G. Muhammad, and M. Alsulaiman, "Physics-informed attention temporal convolutional network for EEG-based motor imagery classification," *IEEE Transactions on Industrial Informatics*, vol. 19, no. 2, pp. 2249–2258, 2023.

[18] Y. Song, Q. Zheng, B. Liu, and X. Gao, "EEG Conformer: Convolutional transformer for EEG decoding and visualization," *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 31, pp. 710–719, 2023.

[19] F. Scholkmann, S. Kleiser, A. J. Metz, R. Zimmermann, J. Mata Pavia, U. Wolf, and M. Wolf, "A review on continuous wave functional near-infrared spectroscopy and imaging instrumentation and methodology," *NeuroImage*, vol. 85, pp. 6–27, 2014.

[20] P. Pandey, J. McLinden, N. Rahimi, C. Kumar, M. Shao, K. M. Spencer, S. Ostadabbas, and Y. Shahriari, "fNIRSNET: A multi-view spatio-temporal convolutional neural network fusion for functional near-infrared spectroscopy-based auditory event classification," *Engineering Applications of Artificial Intelligence*, vol. 137, p. 109256, 2024.
[21] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 5998–6008.
[22] G. H. Glover, "Deconvolution of impulse response in event-related BOLD fMRI," *NeuroImage*, vol. 9, no. 4, pp. 416–429, 1999.
[23] J. Shin, A. von Lühmann, D.-W. Kim, J. Mehnert, H.-J. Hwang, and K.-R. Müller, "Simultaneous acquisition of EEG and NIRS during cognitive tasks for an open access dataset," *Scientific Data*, vol. 5, no. 1, p. 180003, 2018.
[24] A. von Lühmann, H. Wabnitz, T. Sander, and K.-R. Müller, "M3BA: A mobile, modular, multimodal biosignal acquisition architecture for miniaturized EEG-NIRS-based hybrid BCI and monitoring," *IEEE Transactions on Biomedical Engineering*, vol. 64, no. 6, pp. 1199–1210, 2017.
[25] P. Pinti, I. Tachtsidis, A. Hamilton, J. Hirsch, C. Aichelburg, S. Gilbert, and P. W. Burgess, "The present and future use of functional near-infrared spectroscopy (fNIRS) for cognitive neuroscience," *Annals of the New York Academy of Sciences*, vol. 1464, no. 1, pp. 5–29, 2020.
[26] A. M. Chiarelli, P. Croce, A. Merla, and F. Zappasodi, "Deep learning for hybrid EEG-fNIRS brain-computer interface: application to motor imagery classification," *Journal of Neural Engineering*, vol. 15, no. 3, p. 036028, 2018.
[27] Y. Kwak, W.-J. Song, and S.-E. Kim, "FGANet: fNIRS-guided attention network for hybrid EEG-fNIRS brain-computer interfaces," *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 30, pp. 329–339, 2022.
[28] L. Qiu, W. Feng, L. Zhong, X. Song, Z. Ying, and J. Pan, "Mutual information learning-based end-to-end fusion network for hybrid EEG–fNIRS brain–computer interface," *IEEE Transactions on Instrumentation and Measurement*, vol. 74, pp. 1–13, 2025.
[29] Y. Li, L. Zhu, A. Huang, J. Zhang, and P. Yuan, "Multimodal MBC-ATT: cross-modality attentional fusion of EEG-fNIRS for cognitive state decoding," *Frontiers in Human Neuroscience*, vol. 19, p. 1660532, 2025.
[30] K. J. Friston, P. Fletcher, O. Josephs, A. Holmes, M. D. Rugg, and R. Turner, "Event-related fMRI: characterizing differential responses," *NeuroImage*, vol. 7, no. 1, pp. 30–40, 1998.
[31] G. Strangman, J. P. Culver, J. H. Thompson, and D. A. Boas, "A quantitative comparison of simultaneous BOLD fMRI and NIRS recordings during functional brain activation," *NeuroImage*, vol. 17, no. 2, pp. 719–731, 2002.
[32] D. T. Delpy, M. Cope, P. van der Zee, S. Arridge, S. Wray, and J. Wyatt, "Estimation of optical pathlength through tissue from direct time of flight measurement," *Physics in Medicine & Biology*, vol. 33, no. 12, pp. 1433–1442, 1988.
[33] Y. Benjamini and Y. Hochberg, "Controlling the false discovery rate: a practical and powerful approach to multiple testing," *Journal of the Royal Statistical Society: Series B (Methodological)*, vol. 57, no. 1, pp. 289–300, 1995.
[34] J. R. Wolpaw, N. Birbaumer, D. J. McFarland, G. Pfurtscheller, and T. M. Vaughan, "Brain-computer interfaces for communication and control," *Clinical Neurophysiology*, vol. 113, no. 6, pp. 767–791, 2002.
[35] A. Gevins, M. E. Smith, L. McEvoy, and D. Yu, "High-resolution EEG mapping of cortical activation related to working memory: effects of task difficulty, type of processing, and practice," *Cerebral Cortex*, vol. 7, no. 4, pp. 374–385, 1997.
[36] D. J. McFarland, L. A. Miner, T. M. Vaughan, and J. R. Wolpaw, "Mu and beta rhythm topographies during motor imagery and actual movements," *Brain Topography*, vol. 12, no. 3, pp. 177–186, 2000.