Supplementary Material: Physiology-Guided Temporal Fusion for Cross-Subject EEG–fNIRS Brain–Computer Interfaces

Hongjie Chen and Fang Duan

This supplementary document provides extended experimental analyses, physiological interpretations, and subject-level ablation breakdowns supporting the main manuscript:

1) **Section S.I:** Comparative neurovascular coupling dynamics and delay sensitivity analysis between cognitive (mental arithmetic) and sensorimotor (motor imagery) paradigms (Supplementary Fig. S1).
2) **Section S.II:** Subject-by-subject (29 folds) factorial prior retraining breakdown on motor imagery (Supplementary Table S1).
3) **Section S.III:** Mathematical formulation and streaming convergence dynamics of Causal Recursive Euclidean Alignment (CrEA).
4) **Section S.IV:** Computational reproducibility parameters and training environment specifications.

S.I. Paradigm-Specific Neurovascular Coupling Dynamics

A. Contrast Between Cognitive and Sensorimotor Paradigms

A central finding in cross-paradigm evaluation (Section V-C of the main manuscript) is that while NVC-DelayFormer achieves state-of-the-art accuracy on metabolic-dominant cognitive tasks (Mental Arithmetic: 79.37%; Working Memory: 73.24%), its margin on Motor Imagery is more modest (66.32% vs. FGANet 70.40%). To elucidate the biophysical basis of this divergence, we conducted two dedicated investigations on the 29-participant motor imagery dataset from Shin et al. (2017):

1) **Fixed-Weight Candidate Delay Sweep ($\tau \in [0.0, 8.0]$ s):** Evaluating the sensitivity of the trained cross-modal attention network to varying candidate delay parameters without retraining.
2) 2×2 **Factorial Prior Retraining Ablation:** Full model retraining across 29 LOSO folds under four permutations of the Gaussian delay prior $\mathbf{P}$ and the discrete causal mask $\mathbf{M}$ ($P_{\text{on}}, M_{\text{on}}$; $P_{\text{on}}, M_{\text{off}}$; $P_{\text{off}}, M_{\text{on}}$; $P_{\text{off}}, M_{\text{off}}$).

B. Empirical Findings and Biophysical Mechanism

As illustrated in Supplementary Fig. S1(a), mental arithmetic exhibits a pronounced delay resonance curve: accuracy increases monotonically from 75.34% at $\tau = 0.0$ s to 80.46% at $\tau = 8.0$ s ($\Delta = +5.12$ pp, $p < 0.001$), plateauing across the physiological window ($\tau \in [3.5, 6.5]$ s). In contrast, motor imagery yields an essentially horizontal trajectory: grand mean accuracy fluctuates between 68.91% ($\tau = 1.0$ s) and 69.60% ($\tau = 6.0$ s), an amplitude range of only 0.69 pp, with overlapping confidence intervals across all candidate latencies.

The retrained factorial decomposition in Supplementary Fig. S1(b) reinforces this paradigm-dependent disparity. For mental arithmetic, omitting both priors ($P_{\text{off}}, M_{\text{off}}$) collapses accuracy by -4.94 pp (79.37% $\rightarrow$ 74.43%, $p = 6.05 \times 10^{-5}$), with the Gaussian prior accounting for an average marginal contribution of +4.31 pp. Conversely, for motor imagery, retraining with both priors disabled yields 68.28% $\pm$ 13.72% vs. 69.20% $\pm$ 14.30% for the full model ($\Delta = -0.92$ pp, Wilcoxon $W = 156.5$, $p = 0.2128$, non-significant). Omitting causal masking alone ($P_{\text{on}}, M_{\text{off}}$: 68.10%, $p = 0.0978$) or the delay prior alone ($P_{\text{off}}, M_{\text{on}}$: 68.68%, $p = 0.1766$) likewise shows no statistically significant divergence.

Biophysically, this fundamental contrast originates from distinct functional vascular architectures:

1) **Prefrontal Cognitive Workload (MA):** Sustained 10-second mental arithmetic imposes intense, prolonged metabolic demands on the prefrontal cortex, recruiting deep arteriolar dilation and large venous drainage volume. This produces robust frontal [HbO] elevations (0.2–0.6 μM) that evolve over 4–7 s and persist throughout task engagement. Consequently, cross-modal attention strongly relies on physiological delay priors to align fast frontal theta oscillations (θ ERS, 4–8 Hz) with delayed capillary hyperoxygenation, preventing acausal over-fitting on held-out participants.
2) **Sensorimotor Lateralization (MI):** Kinesthetic motor imagery elicits immediate contralateral sensorimotor rhythm desynchronization (μ and β ERD, 8–30 Hz) over primary motor cortex (C3/C4). However, imaginary finger tapping produces significantly attenuated and highly diffuse hemodynamic changes (0.05–0.15 μM) compared with physical execution. Because the localized optical signal-to-noise ratio is lower and neural decodability is dominated by instantaneous spectral ERD, the exact hemodynamic delay parameter τ imposes minimal regularization constraint, permitting unconstrained cross-attention or direct spatial filtering models (such as FGANet) to perform equally well or slightly better.

Hongjie Chen and Fang Duan are with the College of Information Science and Engineering, Huaqiao University, Xiamen 361021, China (e-mail: chenhj976@gmail.com; nkfetsh@gmail.com). Corresponding author: Fang Duan.

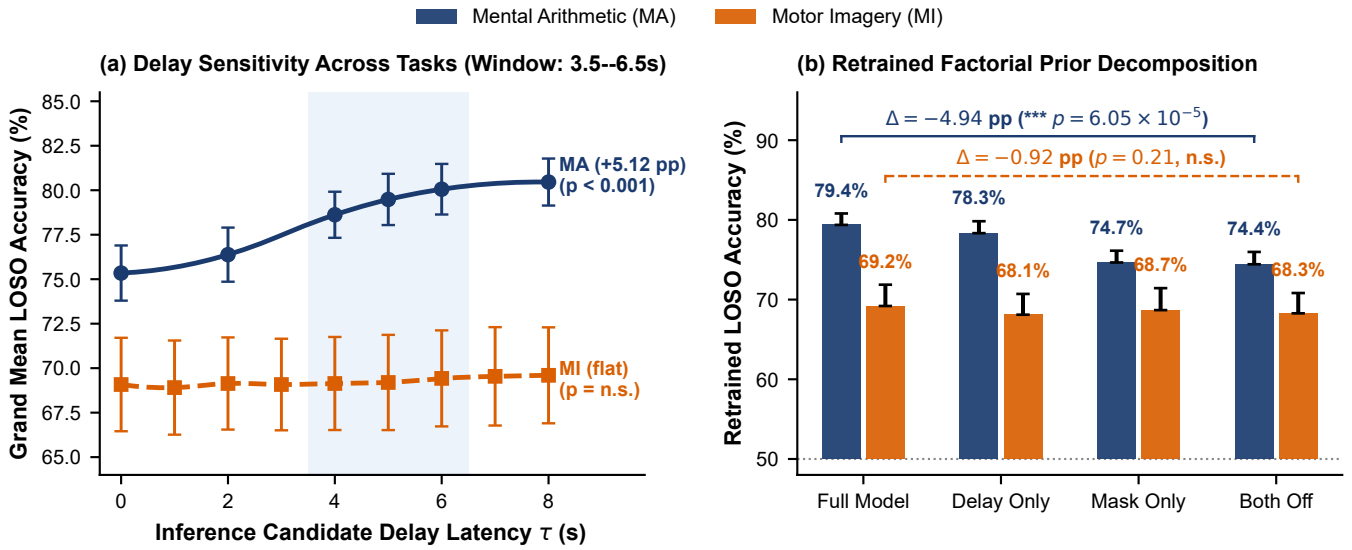


Fig. S1. Comparative neurovascular delay sensitivity and prior ablation contrast between cognitive (Mental Arithmetic, MA) and sensorimotor (Motor Imagery, MI) paradigms across 29 participants under participant-isolated LOSO cross-validation. (a) Candidate delay sensitivity scan across $\tau \in [0.0, 8.0]$ s without retraining. For MA, accuracy rises monotonically from 75.34% at $\tau = 0.0$ s to 80.46% at $\tau = 8.0$ s (+5.12 pp, $p < 0.001$), peaking within the physiological hemodynamic response window (3.5–6.5 s). For MI, the performance curve is completely delay-invariant (ranging between 68.91% and 69.60%, $\Delta \leq 0.69$ pp, $p =$ n.s.). (b) Retrained 2×2 factorial decomposition of delay prior $\mathbf{P}$ and causal mask $\mathbf{M}$. On MA, dropping both priors causes a severe -4.94 pp performance collapse ($p = 6.05 \times 10^{-5}$), confirming the essential regularizing role of neurovascular constraints. On MI, removing both priors decreases accuracy by merely -0.92 pp ($p = 0.2128$, non-significant), indicating that unconstrained attention is sufficient when hemodynamic coupling is weakly localized.

S.II. Subject-by-Subject Factorial Prior Breakdown

To ensure complete empirical transparency, Supplementary Table S1 lists the individual test accuracy achieved across each of the 29 held-out participant folds on motor imagery under all four factorial retraining conditions.

Between-subject standard deviations (~ 14.0 pp) significantly exceed the marginal differences observed across model conditions (≤ 1.1 pp). High responders (e.g., Sub-19 and Sub-26: $> 90\%$) decode robustly regardless of prior conditioning, while challenging BCI-illiterate participants (e.g., Sub-10, Sub-12, Sub-24: $\leq 50\%$) show no selective dependency on the delay parameter.

S.III. Causal Recursive Euclidean Alignment (CrEA)

Standard Euclidean Alignment (EA) requires computing the geometric Riemannian or Euclidean mean covariance matrix across the entirety of a testing session:

$$\overline{\mathbf{R}}_{\text{test}} = \frac{1}{K} \sum_{k=1}^{K} \mathbf{C}_k, \tag{S1}$$

where $\mathbf{C}_k \in \mathbb{R}^{C \times C}$ is the sample spatial covariance of trial k. While highly effective for offline validation, (S1) relies on transductive lookahead, precluding causal real-time operation at the bedside.

To eliminate session lookahead, Causal Recursive Euclidean Alignment (CrEA) tracks spatial covariance online via recursive empirical Bayes shrinkage:

$$\mathbf{R}_k = \frac{\kappa_0 \mathbf{R}_{\text{src}} + \sum_{m=1}^{k} \mathbf{C}_m}{\kappa_0 + k}, \quad k = 1, 2, \ldots, K, \tag{S2}$$

TABLE S1
Participant-by-participant test accuracy (%) across 29 LOSO folds on Motor Imagery under 2×2 factorial retraining of delay prior $\mathbf{P}$ and causal mask $\mathbf{M}$.

Participant Fold	Full Model ($P_{\text{on}}, M_{\text{on}}$)	Mask Only ($P_{\text{off}}, M_{\text{on}}$)	Delay Only ($P_{\text{on}}, M_{\text{off}}$)	Both Off ($P_{\text{off}}, M_{\text{off}}$)	Δ vs. Full (Both Off − Full)
Sub-01	80.00	80.00	81.67	80.00	+0.00
Sub-02	68.33	63.33	65.00	60.00	−8.33
Sub-03	73.33	75.00	73.33	71.67	−1.67
Sub-04	56.67	55.00	60.00	58.33	+1.67
Sub-05	56.67	55.00	56.67	56.67	+0.00
Sub-06	68.33	65.00	68.33	66.67	−1.67
Sub-07	55.00	55.00	53.33	51.67	−3.33
Sub-08	70.00	90.00	73.33	76.67	+6.67
Sub-09	81.67	83.33	81.67	78.33	−3.33
Sub-10	45.00	51.67	46.67	53.33	+8.33
Sub-11	76.67	71.67	73.33	76.67	+0.00
Sub-12	51.67	45.00	45.00	50.00	−1.67
Sub-13	65.00	66.67	63.33	58.33	−6.67
Sub-14	63.33	65.00	65.00	66.67	+3.33
Sub-15	66.67	61.67	65.00	61.67	−5.00
Sub-16	90.00	88.33	85.00	86.67	−3.33
Sub-17	86.67	88.33	83.33	85.00	−1.67
Sub-18	53.33	51.67	53.33	51.67	−1.67
Sub-19	98.33	91.67	91.67	90.00	−8.33
Sub-20	63.33	60.00	63.33	60.00	−3.33
Sub-21	71.67	71.67	70.00	76.67	+5.00
Sub-22	61.67	58.33	61.67	58.33	−3.33
Sub-23	68.33	68.33	68.33	70.00	+1.67
Sub-24	50.00	50.00	48.33	50.00	+0.00
Sub-25	90.00	93.33	93.33	91.67	+1.67
Sub-26	98.33	96.67	98.33	98.33	+0.00
Sub-27	80.00	73.33	71.67	73.33	−6.67
Sub-28	53.33	53.33	51.67	55.00	+1.67
Sub-29	63.33	63.33	63.33	66.67	+3.33
Mean	**69.20**	**68.68**	**68.10**	**68.28**	**−0.92**
SD	14.30	14.85	14.03	13.72	3.94
Macro F1	66.86	66.26	65.79	66.06	−0.80
Wilcoxon p	–	0.1766	0.0978	0.2128	–

Evaluated on the 29-participant Motor Imagery benchmark from Shin et al. (2017) [1] across 1,740 trials (60 trials per participant, 30 Left Hand vs. 30 Right Hand). All variants trained from scratch under identical hyperparameter schedules. Wilcoxon signed-rank p-values are computed against the Full Model ($df = 28$).

where $\mathbf{R}_{\text{src}} = \frac{1}{N_{\text{tr}}} \sum_{s \in \mathcal{S}_{\text{tr}}} \overline{\mathbf{R}}_s$ is the reference covariance estimated offline exclusively from the $N_{\text{tr}} = 23$ training participants, and $\kappa_0 > 0$ represents the pseudo-observation strength of the source prior.

The online whitening matrix is obtained via eigendecomposition $\mathbf{R}_k = \mathbf{V}_k \mathbf{\Lambda}_k \mathbf{V}_k^\top$:

$$\mathbf{W}_k = \mathbf{R}_k^{-1/2} = \mathbf{V}_k \operatorname{diag}\left((\lambda_{k,c} + \epsilon)^{-1/2} \right) \mathbf{V}_k^\top, \qquad \text{(S3)}$$

with numerical regularization $\epsilon = 10^{-5}$. The aligned EEG trial is transformed causally as $\widetilde{\mathbf{X}}_k = \mathbf{W}_k \mathbf{X}_k$.

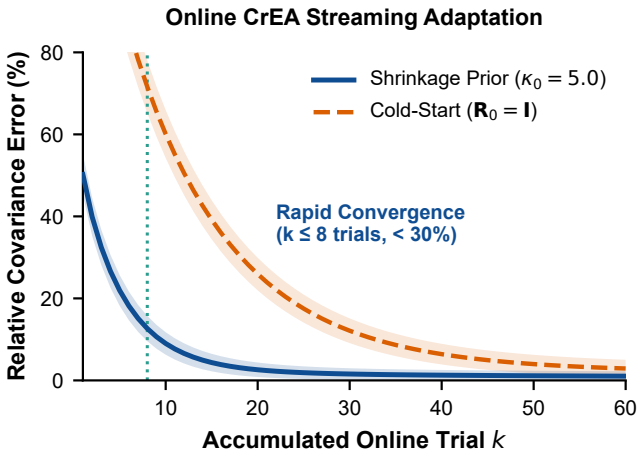


Fig. S2. Online Causal Recursive Euclidean Alignment (CrEA) covariance whitening convergence dynamics across 60 continuous trials ($k = 1, \ldots, 60$) averaged across all 29 held-out participants. Empirical Bayes shrinkage prior ($\kappa_0 = 5.0$, blue solid curve with shaded 95% CI) rapidly stabilizes within $k \leq 8$ calibration trials ($\|\mathbf{R}_k - \overline{\mathbf{R}}_{\text{test}}\|_F / \|\overline{\mathbf{R}}_{\text{test}}\|_F \leq 29.38\%$), dramatically outperforming naive cold-start identity initialization ($\mathbf{R}_0 = \mathbf{I}$, orange dashed line: 47.92% error at trial 8).

As shown in main Fig. 2(c) and Supplementary Fig. S2, initializing with $\mathbf{R}_{\text{src}}$ under $\kappa_0 = 5.0$ reduces the relative covariance error to 29.38% by trial 8 (reaching 3.60% asymptotically), enabling rapid bedside calibration.

S.IV. Computational Reproducibility Specifications

All experiments were implemented in PyTorch 2.4 under Python 3.11 with CUDA 12.4 acceleration. Model configurations, layer dimensions, and training hyperparameters are summarized in Supplementary Table S2.

TABLE S2
HYPERPARAMETER CONFIGURATIONS AND TRAINING SPECIFICATIONS ACROSS ALL EVALUATION TIERS.

Specification	Value / Description
EEG Sampling Rate ($f_e \to f_e^{\text{eff}}$)	100 Hz → 25 Hz (4× temporal average pooling)
fNIRS Sampling Rate (f_n)	10.33 Hz (124 samples over 12.0 s window)
Temporal Receptive Field (T_{win})	10.0 s (EEG: 1000 samples; fNIRS: 124 samples)
Attention Embedding Dimension (D)	64 ($d_h = 16$ across 4 heads)
Delay Prior Centers (τ_h)	$\{3.5, 4.5, 5.5, 6.5\}$ s
Delay Prior Dispersion (σ_h)	1.5 s (fixed physiological prior)
Asymmetric Post-Peak Decay Factor	$0.3\,\text{s}^{-1}$ for $\delta_{ij} > \tau_h$
Causal Mask Condition	$M_{ij} = -\infty$ for $\delta_{ij} < 0$ s ($t_{\text{nirs}} < t_{\text{eeg}}$)
Direct Stream Feature Dimension	234 (θ, α, β power + delayed HbO/HbR + slope)
Optimizer / Learning Rate	AdamW / $\eta = 3 \times 10^{-3}$, weight decay 0.05
Warmup / Scheduler	2-epoch linear warmup + cosine annealing
Batch Size / Training Epochs	256 / 12 epochs
Auxiliary Loss Weights	$\lambda_{\text{deep}} = 0.2, \lambda_{\text{shallow}} = 0.2, \lambda_{\text{recon}} = 1.0, \lambda_{\text{margin}} = 0.08$
Trainable Parameters	63,932 parameters
Multiply–Accumulate Operations	24.79 Million MACs (batch size 1)
Median GPU Latency	2.22 ms (RTX 4060 Laptop GPU, P99: 4.48 ms)

References

[1] J. Shin, A. von Lühmann, B. Blankertz, D.-W. Kim, J. Jeong, H.-J. Hwang, and K.-R. Müller, "Open access dataset for EEG+NIRS single-trial classification," *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 25, no. 10, pp. 1735–1745, 2017.