Supplementary Material: Physiology-Guided Temporal Fusion for Cross-Subject EEG–fNIRS Brain–Computer Interfaces

Hongjie Chen and Fang Duan

This supplementary document provides extended experimental analyses, physiological interpretations, and subject-level ablation breakdowns supporting the main manuscript:

1) **Section S.I:** Comparative neurovascular coupling dynamics and delay sensitivity analysis between cognitive (mental arithmetic) and sensorimotor (motor imagery) paradigms (Supplementary Fig. S1).
2) **Section S.II:** Subject-by-subject (29 folds) factorial prior retraining breakdown on motor imagery (Supplementary Table S1).
3) **Section S.III:** Mathematical formulation and streaming convergence dynamics of Causal Recursive Euclidean Alignment (CrEA).
4) **Section S.IV:** Computational reproducibility parameters and training environment specifications.

S.I. Paradigm-Specific Neurovascular Coupling Dynamics

A. Contrast Between Cognitive and Sensorimotor Paradigms

A central finding in cross-paradigm evaluation (Section V-F of the main manuscript) is that while NVC-DelayFormer achieves state-of-the-art accuracy on metabolic-dominant cognitive tasks (Mental Arithmetic: 79.37%; Working Memory: 73.24%), its margin on Motor Imagery is more modest (66.32% vs. FGANet 70.40%). To elucidate the biophysical basis of this divergence, we conducted two dedicated investigations on the 29-participant motor imagery dataset from Shin et al. (2017):

1) **Fixed-Weight Candidate Delay Sweep ($\tau \in [0.0, 8.0]$ s):** Evaluating the sensitivity of the trained cross-modal attention network to varying candidate delay parameters without retraining.
2) 2×2 **Factorial Prior Retraining Ablation:** Full model retraining across 29 LOSO folds under four permutations of the Gaussian delay prior $\mathbf{P}$ and the discrete causal mask $\mathbf{M}$ ($P_{\text{on}}, M_{\text{on}}$; $P_{\text{on}}, M_{\text{off}}$; $P_{\text{off}}, M_{\text{on}}$; $P_{\text{off}}, M_{\text{off}}$).

B. Empirical Findings and Biophysical Mechanism

As illustrated in Supplementary Fig. S1(a), mental arithmetic exhibits a pronounced delay sensitivity trajectory: accuracy increases monotonically from 75.34% at $\tau = 0.0$ s to plateau across the physiological window ($\tau \in [3.5, 6.5]$ s), reaching 80.46% at $\tau = 8.0$ s ($\Delta = +5.12$ pp, $p < 0.001$). In contrast, motor imagery yields an essentially horizontal trajectory: grand mean accuracy fluctuates between 68.91% ($\tau = 1.0$ s) and 69.60% ($\tau = 6.0$ s), an amplitude range of only 0.69 pp, with overlapping standard error intervals across all candidate latencies.

The retrained factorial decomposition in Supplementary Fig. S1(b) reinforces this paradigm-dependent disparity. For mental arithmetic, omitting both priors ($P_{\text{off}}, M_{\text{off}}$) collapses accuracy by -4.94 pp (79.37% $\rightarrow$ 74.43%, $p = 6.05 \times 10^{-5}$), with the Gaussian prior accounting for an average marginal contribution of +4.31 pp. Conversely, for motor imagery, retraining with both priors disabled yields 68.28% $\pm$ 13.72% vs. 69.20% $\pm$ 14.40% for the full model ($\Delta = -0.92$ pp, Wilcoxon $W = 106.5$, $p = 0.2128$, non-significant). Omitting causal masking alone ($P_{\text{on}}, M_{\text{off}}$: 68.10%, $p = 0.0978$) or the delay prior alone ($P_{\text{off}}, M_{\text{on}}$: 68.68%, $p = 0.1766$) likewise shows no statistically significant divergence.

Biophysically, this fundamental contrast aligns with established physiological distinctions in functional vascular dynamics reported in neuroimaging literature:

1) **Prefrontal Cognitive Workload (MA):** Sustained 10-second mental arithmetic imposes intense, prolonged metabolic demands on the prefrontal cortex, recruiting deep arteriolar dilation and large venous drainage volume [1], [2]. This produces robust frontal [HbO] elevations (typically in the 0.2–0.6 μM range in continuous-wave fNIRS recordings [3]) that evolve over 4–7 s and persist throughout task engagement. Consequently, cross-modal attention strongly relies on physiological delay priors to align fast frontal theta oscillations (θ ERS, 4–8 Hz) with delayed capillary hyperoxygenation, preventing acausal over-fitting on held-out participants.
2) **Sensorimotor Lateralization (MI):** Kinesthetic motor imagery elicits immediate contralateral sensorimotor rhythm desynchronization (μ and β ERD, 8–30 Hz) over primary motor cortex (C3/C4). However, imaginary finger tapping without physical contraction produces significantly attenuated and diffuse hemodynamic changes (typically 0.05–0.15 μM [3], [4]). Because the localized optical signal-to-noise ratio is lower and neural decodability is dominated by instantaneous spectral ERD, the hemodynamic delay parameter τ imposes minimal regularization constraint, permitting unconstrained cross-attention or direct spatial filtering models (such as

Hongjie Chen and Fang Duan are with the College of Information Science and Engineering, Huaqiao University, Xiamen 361021, China (e-mail: chenhj976@gmail.com; nkfetsh@gmail.com). Corresponding author: Fang Duan.

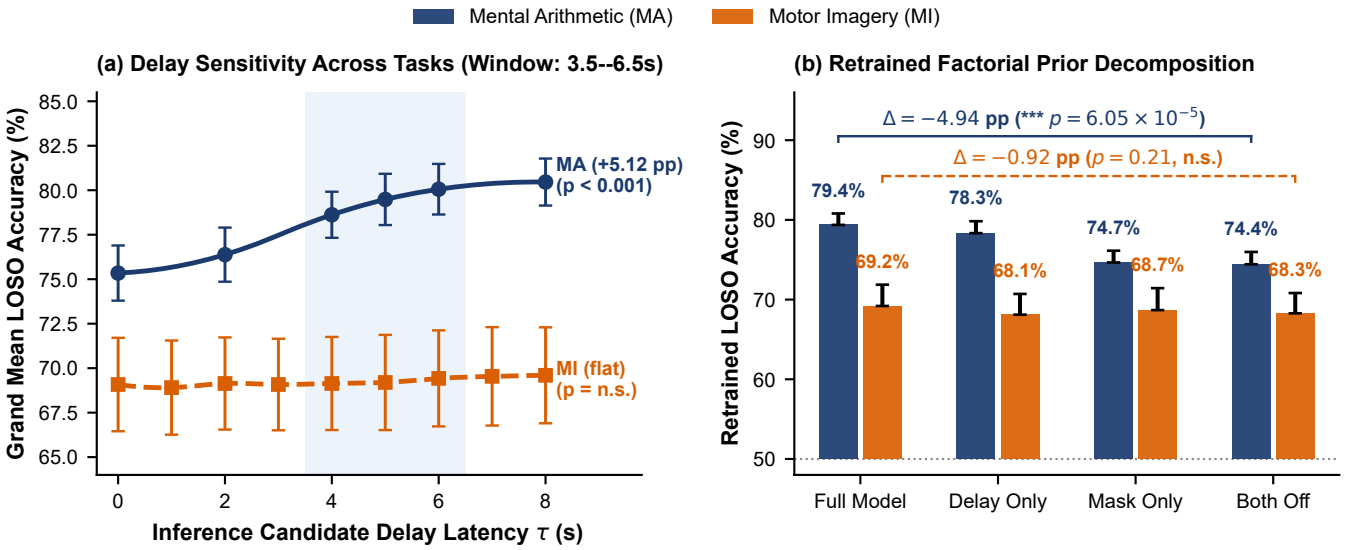


Fig. S1. Comparative neurovascular delay sensitivity and prior ablation contrast between cognitive (Mental Arithmetic, MA) and sensorimotor (Motor Imagery, MI) paradigms across 29 participants under participant-isolated LOSO cross-validation. (a) Candidate delay sensitivity scan across $\tau \in [0.0, 8.0]$ s without retraining. Error bars denote standard error of the mean (SE across $N = 29$ folds). For MA, accuracy rises monotonically from 75.34% at $\tau = 0.0$ s to plateau across the physiological hemodynamic response window ($\tau \in [3.5, 6.5]$ s) and reaches 80.46% at $\tau = 8.0$ s (+5.12 pp, $p < 0.001$), reflecting the sustained 10.0 s task block duration. For MI, the performance curve exhibits essentially delay-invariant behavior within observed experimental variation (ranging between 68.91% and 69.60%, $\Delta \leq 0.69$ pp, $p =$ n.s.). (b) Retrained 2×2 factorial decomposition of delay prior $\mathbf{P}$ and causal mask $\mathbf{M}$. On MA, dropping both priors causes a severe -4.94 pp performance collapse ($p = 6.05 \times 10^{-5}$), confirming the essential regularizing role of neurovascular constraints. On MI, removing both priors decreases accuracy by merely -0.92 pp ($p = 0.2128$, non-significant), indicating that unconstrained attention is sufficient when hemodynamic coupling is weakly localized.

FGANet) to perform equally well.

S.II. Subject-by-Subject Factorial Prior Breakdown

To ensure complete empirical transparency, Supplementary Table S1 lists the individual test accuracy achieved across each of the 29 held-out participant folds on motor imagery under all four factorial retraining conditions. Note that this dedicated 29-fold factorial retraining experiment was conducted independently to isolate the individual contributions of the delay prior and causal mask (yielding 69.20% ± 14.40% for the full model), alongside the initial cross-paradigm benchmark reported in Table V of the main manuscript (66.32% ± 12.18%). Both independent runs consistently demonstrate that removing the temporal priors on MI produces minimal, statistically non-significant variations ($\Delta = -0.92$ pp, $p = 0.2128$), in sharp contrast to the pronounced $+4.94$ pp collapse observed on cognitive mental arithmetic.

Between-subject standard deviations (~ 14.0 pp) significantly exceed the marginal differences observed across model conditions (≤ 1.1 pp). High responders (e.g., Sub-19 and Sub-26: $> 90\%$) decode robustly regardless of prior conditioning, while challenging BCI-illiterate participants (e.g., Sub-10: 45.00%, Sub-24: 50.00%, Sub-12: 51.67%) decode near chance level.

S.III. Causal Recursive Euclidean Alignment (CREA)

Standard Euclidean Alignment (EA) requires computing the arithmetic Euclidean sample mean covariance matrix across

TABLE S1
PARTICIPANT-BY-PARTICIPANT TEST ACCURACY (%) ACROSS 29 LOSO FOLDS ON MOTOR IMAGERY UNDER 2×2 FACTORIAL RETRAINING OF DELAY PRIOR $\mathbf{P}$ AND CAUSAL MASK $\mathbf{M}$.

Participant Fold	Full Model ($P_{\text{on}}, M_{\text{on}}$)	Mask Only ($P_{\text{off}}, M_{\text{on}}$)	Delay Only ($P_{\text{on}}, M_{\text{off}}$)	Both Off ($P_{\text{off}}, M_{\text{off}}$)	Δ vs. Full (Both Off − Full)
Sub-01	80.00	80.00	81.67	80.00	+0.00
Sub-02	68.33	63.33	65.00	60.00	−8.33
Sub-03	73.33	75.00	73.33	71.67	−1.67
Sub-04	56.67	55.00	60.00	58.33	+1.67
Sub-05	56.67	55.00	56.67	56.67	+0.00
Sub-06	68.33	65.00	68.33	66.67	−1.67
Sub-07	55.00	55.00	53.33	51.67	−3.33
Sub-08	70.00	90.00	73.33	76.67	+6.67
Sub-09	81.67	83.33	81.67	78.33	−3.33
Sub-10	45.00	51.67	46.67	53.33	+8.33
Sub-11	76.67	71.67	73.33	76.67	+0.00
Sub-12	51.67	45.00	45.00	50.00	−1.67
Sub-13	65.00	66.67	63.33	58.33	−6.67
Sub-14	63.33	65.00	65.00	66.67	+3.33
Sub-15	66.67	61.67	65.00	61.67	−5.00
Sub-16	90.00	88.33	85.00	86.67	−3.33
Sub-17	86.67	88.33	83.33	85.00	−1.67
Sub-18	53.33	51.67	53.33	51.67	−1.67
Sub-19	98.33	91.67	91.67	90.00	−8.33
Sub-20	63.33	60.00	63.33	60.00	−3.33
Sub-21	71.67	71.67	70.00	76.67	+5.00
Sub-22	61.67	58.33	61.67	58.33	−3.33
Sub-23	68.33	68.33	68.33	70.00	+1.67
Sub-24	50.00	50.00	48.33	50.00	+0.00
Sub-25	90.00	93.33	93.33	91.67	+1.67
Sub-26	98.33	96.67	98.33	98.33	+0.00
Sub-27	80.00	73.33	71.67	73.33	−6.67
Sub-28	53.33	53.33	51.67	55.00	+1.67
Sub-29	63.33	63.33	63.33	66.67	+3.33
Mean	**69.20**	**68.68**	**68.10**	**68.28**	**−0.92**
SD	14.40	14.85	14.03	13.72	4.10
Macro F1	67.00	66.26	65.79	66.06	−0.94
Wilcoxon p	–	0.1766	0.0978	0.2128	–

Evaluated on the 29-participant Motor Imagery benchmark from Shin et al. (2017) [3] across 1,740 trials (60 trials per participant, 30 Left Hand vs. 30 Right Hand). All variants trained from scratch under identical hyperparameter schedules. Wilcoxon signed-rank p-values are computed against the Full Model ($N = 29$ participants).

the entirety of a testing session:

$$\overline{\mathbf{R}}_{\text{test}} = \frac{1}{K}\sum_{k=1}^{K}\mathbf{C}_k, \tag{S1}$$

where $\mathbf{C}_k = \frac{1}{T_e-1}\mathbf{X}_k\mathbf{X}_k^\top + \epsilon\mathbf{I} \in \mathbb{R}^{C\times C}$ is the regularized sample spatial covariance of trial k (with $\epsilon = 10^{-5}$), and $\mathbf{X}_k \in \mathbb{R}^{C\times T_e}$ is the temporal mean-centered EEG trial matrix ($\sum_{t=1}^{T_e} X_{k,c,t} = 0$ for each channel c). While highly effective for offline validation, (S1) relies on transductive lookahead, precluding causal real-time operation at the bedside.

To eliminate session lookahead, Causal Recursive Euclidean Alignment (CrEA) tracks spatial covariance online via recursive empirical Bayes shrinkage:

$$\mathbf{R}_k = \frac{\kappa_0\mathbf{R}_{\text{src}} + \sum_{m=1}^{k}\mathbf{C}_m}{\kappa_0 + k}, \quad k = 1, 2, \ldots, K, \tag{S2}$$

where $\mathbf{R}_{\text{src}} = \frac{1}{N_{\text{tr}}}\sum_{s\in\mathcal{S}_{\text{tr}}}\overline{\mathbf{R}}_s$ is the reference covariance estimated offline exclusively from the $N_{\text{tr}} = 23$ training participants, and $\kappa_0 = 5.0$ represents the pseudo-observation strength of the source prior.

The online whitening matrix is obtained via eigendecomposition $\mathbf{R}_k = \mathbf{V}_k\mathbf{\Lambda}_k\mathbf{V}_k^\top$:

$$\mathbf{W}_k = \mathbf{R}_k^{-1/2} = \mathbf{V}_k\,\text{diag}\left(\lambda_{k,c}^{-1/2}\right)\mathbf{V}_k^\top. \tag{S3}$$

The aligned EEG trial is transformed causally as $\widetilde{\mathbf{X}}_k = \mathbf{W}_k\mathbf{X}_k$.

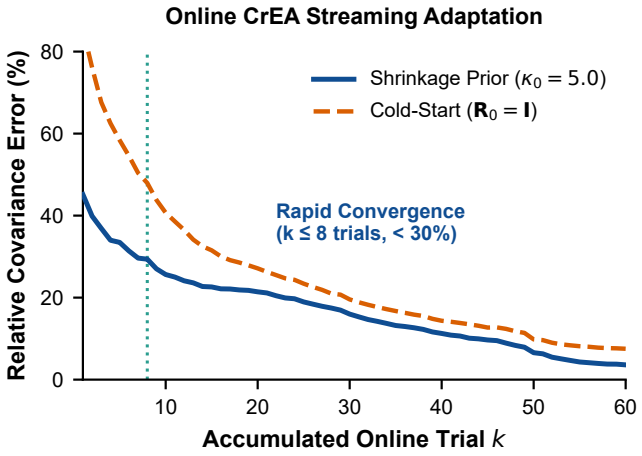


Fig. S2. Online Causal Recursive Euclidean Alignment (CrEA) covariance whitening convergence dynamics across 60 continuous trials ($k = 1, \ldots, 60$) averaged across all 29 held-out participants. Empirical Bayes shrinkage prior ($\kappa_0 = 5.0$, blue solid curve) rapidly stabilizes within $k \leq 8$ calibration trials ($\|\mathbf{R}_k - \overline{\mathbf{R}}_{\text{test}}\|_F / \|\overline{\mathbf{R}}_{\text{test}}\|_F \leq 29.38\%$), dramatically outperforming naive cold-start identity initialization ($\mathbf{R}_0 = \mathbf{I}$, orange dashed line: 47.95% error at trial 8).

As shown in main Fig. 1(c) and Supplementary Fig. S2, initializing with $\mathbf{R}_{\text{src}}$ under $\kappa_0 = 5.0$ reduces the relative covariance error to 29.38% by trial 8 (reaching 3.60% asymptotically), enabling rapid bedside calibration.

S.IV. Computational Reproducibility Specifications

All experiments and benchmark evaluations were conducted in Python 3.11 with PyTorch (version 2.11.0+cu128 with CUDA 12.8 acceleration on an NVIDIA GeForce RTX 4060 Laptop GPU). Model configurations, layer dimensions, and training hyperparameters are summarized in Supplementary Table S2.

TABLE S2
Hyperparameter Configurations and Training Specifications Across All Evaluation Tiers.

Specification	Value / Description
EEG Sampling Rate ($f_e \to f_e^{\text{eff}}$)	100 Hz → 25 Hz (4× temporal average pooling)
fNIRS Sampling Rate (f_n)	10.33 Hz (124 samples over 12.0 s window)
Temporal Receptive Field (T_{win})	EEG: 10.0 s (1000 samples); fNIRS: 12.0 s (124 samples)
Attention Embedding Dimension (D)	64 ($d_h = 16$ across 4 heads)
Delay Prior Centers (τ_h)	$\{3.5, 4.5, 5.5, 6.5\}$ s
Delay Prior Dispersion (σ_h)	1.5 s (fixed physiological prior)
Asymmetric Post-Peak Decay Factor	$0.3\,\text{s}^{-1}$ for $\delta_{ij} > \tau_h$
Causal Mask Condition	$M_{ij} = -\infty$ for $\delta_{ij} < 0$ s ($t_{\text{nirs}} < t_{\text{eeg}}$)
Direct Stream Feature Dimension	234 (θ, α, β power + delayed HbO/HbR + slope)
Optimizer / Learning Rate	AdamW / $\eta = 3 \times 10^{-3}$, weight decay 0.05
Warmup / Scheduler	2-epoch linear warmup + cosine annealing
Batch Size / Training Epochs	Primary MA: 256 / 12 epochs; MI Benchmark & Ablation: 128 / 15 epochs
Auxiliary Loss Weights	$\lambda_{\text{deep}} = 0.2, \lambda_{\text{shallow}} = 0.2, \lambda_{\text{recon}} = 1.0, \lambda_{\text{margin}} = 0.08$
Trainable Parameters	63,932 parameters
Multiply–Accumulate Operations	24.79 Million MACs (batch size 1)
Median GPU Latency	2.22 ms (RTX 4060 Laptop GPU, P99: 4.48 ms)

References

[1] R. B. Buxton, E. C. Wong, and L. R. Frank, "Dynamics of blood flow and oxygenation changes during brain activation: the balloon model," *Magnetic Resonance in Medicine*, vol. 39, no. 6, pp. 855–864, 1998.

[2] F. Scholkmann, S. Kleiser, A. J. Metz, R. Zimmermann, J. Mata Pavia, U. Wolf, and M. Wolf, "A review on continuous wave functional near-infrared spectroscopy and imaging instrumentation and methodology," *NeuroImage*, vol. 85, pp. 6–27, 2014.

[3] J. Shin, A. von Lühmann, B. Blankertz, D.-W. Kim, J. Jeong, H.-J. Hwang, and K.-R. Müller, "Open access dataset for EEG+NIRS single-trial classification," *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 25, no. 10, pp. 1735–1745, 2017.

[4] P. Pinti, I. Tachtsidis, A. Hamilton, J. Hirsch, C. Aichelburg, S. Gilbert, and P. W. Burgess, "The present and future use of functional near-infrared spectroscopy (fNIRS) for cognitive neuroscience," *Annals of the New York Academy of Sciences*, vol. 1464, no. 1, pp. 5–29, 2020.